

Computing top- k Closeness Centrality Faster in Unweighted Graphs

ELISABETTA BERGAMINI, Karlsruhe Institute of Technology

MICHELE BORASSI, IMT Institute for Advanced Studies Lucca

PIERLUIGI CRESCENZI, Università degli Studi di Firenze and Université de Paris

ANDREA MARINO, Università degli Studi di Firenze

HENNING MEYERHENKE, Karlsruhe Institute of Technology and Humboldt-Universität zu Berlin

Given a connected graph $G = (V, E)$, where V denotes the set of nodes and E the set of edges of the graph, the length (that is, the number of edges) of the shortest path between two nodes v and w is denoted by $d(v, w)$. The closeness centrality of a vertex v is then defined as $\frac{n-1}{\sum_{w \in V} d(v, w)}$, where $n = |V|$. This measure is widely used in the analysis of real-world complex networks, and the problem of selecting the k most central vertices has been deeply analyzed in the last decade. However, this problem is computationally not easy, especially for large networks: in the first part of the article, we prove that it is not solvable in time $O(|E|^{2-\epsilon})$ on directed graphs, for any constant $\epsilon > 0$, under reasonable complexity assumptions. Furthermore, we propose a new algorithm for selecting the k most central nodes in a graph: we experimentally show that this algorithm improves significantly both the textbook algorithm, which is based on computing the distance between all pairs of vertices, and the state of the art. For example, we are able to compute the top k nodes in few dozens of seconds in real-world networks with millions of nodes and edges. Finally, as a case study, we compute the 10 most central actors in the Internet Movie Database (IMDB) collaboration network, where two actors are linked if they played together in a movie, and in the Wikipedia citation network, which contains a directed edge from a page p to a page q if p contains a link to q .

CCS Concepts: • **Human-centered computing** \rightarrow **Social network analysis**; • **Mathematics of computing** \rightarrow **Graph algorithms**;

Additional Key Words and Phrases: Centrality, closeness, complex networks

E. Bergamini's and H. Meyerhenke's work is partially supported by German Research Foundation (DFG) grant ME 3619/3-2 within the Priority Programme 1736 *Algorithms for Big Data*. P. Crescenzi's and A. Marino's work is partially supported by the Italian Ministry of Education, University, and Research (MIUR) under PRIN 2012C4E3KT national research project AMANDA—Algorithmics for MAssive and Networked DATa. This paper is an extended and improved version of the conference paper (Bergamini et al. 2016). A preliminary study appeared in (Borassi et al. 2015b).

Authors' addresses: E. Bergamini, Institute of Theoretical Informatics, Karlsruhe Institute of Technology (KIT), Karlsruhe, Germany; M. Borassi, IMT Institute for Advanced Studies Lucca, Italy; P. Crescenzi, Université de Paris, IRIF, Bâtiment Sophie Germain, 8 Place Aurélie Nemours, 75205 Paris Cedex 13, France; email: pierluigi.crescenzi@irif.fr; A. Marino, Dipartimento di Statistica, Informatica, Applicazioni 'G. Parenti' (DiSIA), Università degli Studi di Firenze, Viale Morgagni 59, 50134 Firenze, Italy; email: andrea.marino@unifi.it; H. Meyerhenke, Department of Computer Science, Humboldt-Universität zu Berlin, Unter den Linden 6, 10099 Berlin, Germany; email: meyerhenke@hu-berlin.de.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2019 Association for Computing Machinery.

1556-4681/2019/09-ART53 \$15.00

<https://doi.org/10.1145/3344719>

ACM Reference format:

Elisabetta Bergamini, Michele Borassi, Pierluigi Crescenzi, Andrea Marino, and Henning Meyerhenke. 2019. Computing top- k Closeness Centrality Faster in Unweighted Graphs. *ACM Trans. Knowl. Discov. Data* 13, 5, Article 53 (September 2019), 40 pages.
<https://doi.org/10.1145/3344719>

1 INTRODUCTION

The problem of identifying the most central nodes in a network is a fundamental question that has been asked many times in a plethora of research areas, such as biology, computer science, sociology, and psychology. Because of the importance of this question, dozens of centrality measures have been introduced in the literature (for a recent survey, see (Boldi and Vigna 2014)). Among these measures, closeness centrality is certainly one of the oldest and of the most widely used (Bavelas 1950): almost all books dealing with network analysis discuss it (for example, (Newman 2010)), and almost all existing network analysis libraries implement algorithms to compute it.

Given a connected graph $G = (V, E)$, where V denotes the set of nodes and E the set of edges of the graph, the length (that is, the number of edges) of the shortest path between two nodes v and w is denoted by $d(v, w)$. The closeness centrality of a vertex v is then defined as $c(v) = \frac{n-1}{\sum_{w \in V} d(v, w)}$, where $n = |V|$. The idea behind this definition is that a central node should be very efficient in spreading information to all other nodes: for this reason, a node is central if the average number of links needed to reach another node is small. If the graph is not (strongly) connected, researchers have proposed various ways to extend this definition: for the sake of simplicity, we focus on Lin's index, because it coincides with closeness centrality in the connected case and because it is quite established in the literature (Lin 1976; Wasserman and Faust 1994; Boldi and Vigna 2013, 2014; Olsen et al. 2014). However, our algorithms can be adapted very easily to compute other possible generalizations, such as harmonic centrality (Marchiori and Latora 2000) and exponential centrality (Wang and Tang 2014) (see Section 2 for more details).

In order to compute the k vertices with largest closeness, the textbook algorithm computes $c(v)$ for each v and returns the k largest found values. The main bottleneck of this approach is the computation of $d(v, w)$ for each pair of vertices v and w (that is, solving the All Pairs Shortest Paths or APSP problem). This can be done in two ways: either by using fast matrix multiplication (for which the currently best algorithms require time $O(n^{2.373})$ (Williams 2012; Gall 2014)) or by performing a breadth-first search (in short, BFS) from each vertex $v \in V$, in time $O(mn)$, where $m = |E|$. Usually, BFS is preferred because the approaches based on fast matrix multiplication present a reduced numerical stability and large constants hidden in the big O notation (Gall 2012) (at least the ones that have been implemented so far). Also, real-world networks are usually sparse, that is, m is not much bigger than n . However, even the BFS-based approach is too time-consuming if the input graph is very big (with millions of nodes and hundreds of millions of edges).

Our first result proves that, in the worst case, the BFS-based approach cannot be improved, under reasonable complexity assumptions. In particular, we show that, assuming the Strong Exponential Time Hypothesis (SETH, (Impagliazzo et al. 2001)), the most central vertex cannot be computed in $O(n^{2-\epsilon})$ on sparse graphs, for any $\epsilon > 0$. Consequently, for any $\epsilon > 0$, there is no algorithm to compute the most central vertex in $O(m^{2-\epsilon})$ on general graphs, since otherwise such an algorithm would compute the most central vertex in $O(n^{2-\epsilon})$ on sparse graphs.

Knowing that the BFS-based algorithm cannot be improved in the worst case, in the second part of the article we provide a new exact algorithm that performs much better on real-world networks, making it possible to compute the k most central vertices in networks with millions of nodes

and hundreds of millions of edges. The new approach combines the BFS-based algorithm with a pruning technique: during the algorithm, we compute and update upper bounds on the closeness of all the nodes, and we exclude a node v from the computation as soon as its upper bound is “small enough,” that is, we are sure that v does not belong to the top k nodes. We propose two different strategies to set the initial bounds, and two different strategies to update the bounds during the computation: this means that our algorithm comes in four different variations. The experimental results show that different variations perform well on different kinds of networks, and the best variation of our algorithm drastically outperforms both a probabilistic approach (Okamoto et al. 2008), and the best exact algorithm available until now (Olsen et al. 2014). We have computed for the first time the 10 most central nodes (with respect to the closeness measure) in networks with millions of nodes and hundreds of millions of edges, and we have done so in very little time. A significant example is the wiki-Talk network, which was also used in (Sariyüce et al. 2013), where the authors propose an algorithm to update closeness centralities after edge additions or deletions. Our performance is about 30,000 times better than the performance of the textbook algorithm: if only the most central node is needed, we can recompute it from scratch more than 150 times faster than the geometric average update time in (Sariyüce et al. 2013). Moreover, our approach is not only very efficient, but it is also very easy to code, making it a very good candidate to be implemented in existing graph libraries. We provide an implementation of it in NetworKit (Staudt et al. 2016) and of one of its variations in Sagemath (Csárdi and Nepusz 2006). We sketch the main ideas of the algorithm in Section 4, and we provide all details in Sections 5–8. We experimentally evaluate the efficiency of the new algorithm in Section 9.

Also, our approach can be easily extended to any centrality measure in the form $c(v) = \sum_{w \neq v} f(d(v, w))$, where f is a decreasing function. Apart from Lin’s index, almost all the approaches that try to generalize closeness centrality to disconnected graphs fall under this category. The most popular among these measures is harmonic centrality (Marchiori and Latora 2000), defined as $h(v) = \sum_{w \neq v} \frac{1}{d(v, w)}$. For the sake of completeness, in Section 9, we show that our algorithm performs well also for this measure.

In the last part of the article (Sections 10 and 11), we consider two case studies: the actor collaboration network (1,797,446 vertices, 72,880,156 edges) and the Wikipedia citation network (4,229,697 vertices, 102,165,832 edges). In the actor collaboration network, we analyze the evolution of the 10 most central vertices, considering snapshots taken every 5 years between 1940 and 2014. The computation was performed in little more than 45 minutes. In the Wikipedia case study, we consider both the standard citation network, that contains a directed edge (p, q) if p contains a link to q , and the reversed network, that contains a directed edge (p, q) if q contains a link to p . For most of these graphs, we are able to compute the 10 most central pages in a few minutes, making them available for further analyses.

1.1 Related Work

Closeness is a “traditional” definition of centrality, and consequently it was not “designed with scalability in mind,” as stated in (Kang et al. 2011). Also in (Chen et al. 2012), it is said that closeness centrality can “identify influential nodes,” but it is “incapable to be applied in large-scale networks due to the computational complexity.” The simplest solution considered was to define different measures that might be related to closeness centrality (Kang et al. 2011).

Hardness results. A different line of research has tried to develop more efficient algorithms, or lower bounds for the complexity of this problem. In particular, in (Borassi et al. 2015a) it is proved that finding the least closeness-central vertex is not subquadratic-time solvable, unless SETH is false. In the same line, it is proved in (Abboud et al. 2016) that finding the most central vertex

is not solvable in $O(m^{2-\epsilon})$, assuming the Hitting Set conjecture. This conjecture is very recent, and there are not strong evidences that it holds, apart from its similarity to the Orthogonal Vector conjecture. Conversely, the Orthogonal Vector conjecture is more established: it is implied both by the Hitting Set conjecture (Abboud et al. 2016), and by SETH (Williams 2005), which is a widely used assumption in the context of polynomial-time reductions (Impagliazzo et al. 2001; Williams 2005; Williams and Williams 2010; Pătraşcu and Williams 2010; Roditty and Williams 2013; Abboud et al. 2014; Abboud and Williams 2014; Abboud et al. 2015; Borassi et al. 2015a; Abboud et al. 2016; Borassi 2016). Similar hardness results were also proved in the dense weighted context (Abboud et al. 2015), by linking the complexity of centrality measures to the complexity of computing the All Pairs Shortest Paths.

Approximation Algorithms. In order to deal with the above hardness results, it is possible to design approximation algorithms: the simplest approach samples the distance between a node v and l other nodes w , and returns the average of all values $d(v, w)$ found (Eppstein and Wang 2004). The time complexity is $O(lm)$, to obtain an approximation $\tilde{c}(v)$ of the centrality of each node v such that $\mathbb{P}(|\frac{1}{\tilde{c}(v)} - \frac{1}{c(v)}| \geq \epsilon D) \leq 2e^{-\Omega(l\epsilon^2)}$, where D is the diameter of the graph (the diameter is the maximum distance between any two connected nodes). A more refined approximation algorithm is provided in (Cohen et al. 2014), which combines the sampling approach with a 3-approximation algorithm: this algorithm has running time $O(lm)$, and it provides an estimate $\tilde{c}(v)$ of the centrality of each node v such that $\mathbb{P}(|\frac{1}{\tilde{c}(v)} - \frac{1}{c(v)}| \geq \frac{\epsilon}{c(v)}) \leq 2e^{-\Omega(l\epsilon^3)}$ (note that, differently from the previous algorithm, this algorithm provides a guarantee on the relative error). The most recent result by Chechik et al. (2015) allows to approximate closeness centrality with a coefficient of variation of ϵ using $O(\epsilon^{-2})$ single-source shortest path (SSSP) computations. Alternatively, one can make the probability that the maximum relative error exceeds ϵ polynomially small by using $O(\epsilon^{-2} \log n)$ SSSP computations.

However, these approximation algorithms have not been specifically designed for ranking nodes according to their closeness centrality, and turning them into a trustable top- k algorithm can be a challenging problem. Indeed, observe that, in many real-world cases, we work with so-called *small-world* networks, having a low diameter. Hence, in a typical graph, the average distance between v and a random node w is between 1 and 10. This implies that most of the n values $\frac{1}{c(v)}$ lie in this range, and that, in order to obtain a reliable ranking, we need the error to be close to $\epsilon = \frac{10}{n}$, which might be very small in the case of the vast majority of real-world networks. As an example, performing $O(\epsilon^{-2})$ SSSPs as in (Chechik et al. 2015) would then require $O(\frac{m}{\epsilon^2}) = O(mn^2)$ time in the unweighted case, which is impractical for large graphs. In the absence of theoretical results, it is, however, worth noting that, as a side effect of our new algorithm, we can now quickly certify in practice, even in the case of very large graphs, how good the ranking produced by these approximation algorithms is. For example, if we run this algorithm on our dataset with the same number of iterations as our algorithm, the relative error guaranteed on the centrality of all the nodes is large (usually, above 50% for $k = 100$), because the algorithm is not tailored to the top- k computation. However, with our algorithm one can show that the ranking obtained is very close to the correct one (usually, more than 95 of the 100 most central nodes according to (Chechik et al. 2015) are actually in the top-100).¹ A theoretical justification of this behavior is, in our opinion, a very interesting open problem.

Finally, an approximation algorithm was proposed in (Okamoto et al. 2008), where the sampling technique developed in (Eppstein and Wang 2004) was used to actually compute the top k vertices:

¹Indeed, we obtained a similar experimental result while dealing with the simpler heuristics consisting in choosing as the sample a set of highest degree nodes slightly larger than the sample chosen by the algorithm in (Chechik et al. 2015).

the result is not exact, but it is exact with high probability. The authors proved that the time complexity of their algorithm is $O(mn^{\frac{2}{3}} \log n)$, under the rather strong assumption that closeness centralities are uniformly distributed between 0 and the diameter D (in the worst case, the time complexity of this algorithm is $O(mn)$).

Heuristics. Other approaches have tried to develop incremental algorithms that might be more suited to real-world networks. For instance, in (Lim et al. 2011), the authors develop heuristics to determine the k most central vertices in an online fashion, i.e., by examining only a small fraction of the network nodes. Furthermore, in (Sariyüce et al. 2013), the authors consider the problem of updating the closeness centrality of all nodes after edge insertions or deletions: in some cases, the time needed for the update could be orders of magnitude smaller than the time needed to recompute all centralities from scratch.

Finally, some works have tried to exploit properties of real-world networks in order to find more efficient algorithms. In (Le Merrer et al. 2014), the authors develop a heuristic to compute the k most central vertices according to different measures. The basic idea is to identify central nodes according to a simple centrality measure (for instance, degree of nodes), and then to inspect a small set of central nodes according to this measure, hoping it contains the top k vertices according to the “complex” measure. The last approach (Olsen et al. 2014), proposed by Olsen et al., tries to exploit the properties of real-world networks in order to develop exact algorithms with worst case complexity $O(mn)$, but performing much better in practice. As far as we know, this is the only exact algorithm that is able to efficiently compute the k most central vertices in networks with up to 1 million nodes, before this work.

Software libraries. Despite this huge amount of research, graph libraries still use the textbook algorithm: among them, Boost Graph Library (Hagberg et al. 2008), igraph (Stein and Joyner 2005), and NetworkX (Siek et al. 2001). This is due to the fact that efficient available exact algorithms for top- k closeness centrality, like (Olsen et al. 2014), are relatively recent and make use of several other non-trivial routines. An implementation of the algorithms presented in this article are now included in the publicly available libraries NetworKit (Staudt et al. 2016)² and Sagemath (Csárdi and Nepusz 2006): in particular, the NetworKit implementation has been used for obtaining our experimental results.

2 PRELIMINARIES

We assume the reader to be familiar with the basic notions of graph theory (see, for example, (Cormen et al. 2009)). Our algorithmic results apply both to undirected and directed graphs. We will make clear in the respective context where results apply to only one of the two. For example, the hardness results in Section 3 apply to directed graphs only. All the notations and definitions used throughout this article are summarized in Table 1 (in any case, all notations are also defined in the text). Here, let us only define precisely the closeness centrality of a vertex v . In a connected graph, the fairness of a node v in a graph $G = (V, E)$ is $f(v) = \frac{\sum_{w \in V} d(v, w)}{n-1}$, and the closeness centrality of v is $\frac{1}{f(v)}$. In the disconnected case, the most natural generalization would be $f(v) = \frac{\sum_{w \in R(v)} d(v, w)}{r(v)-1}$, and $c(v) = \frac{1}{f(v)}$, where $R(v)$ is the set of vertices reachable from v , and $r(v) = |R(v)|$. However, this definition does not capture our intuitive notion of centrality: indeed, if v has only one neighbor w at distance 1, and w has out-degree 0, then v becomes very central according to this measure, even if v is intuitively peripheral. For this reason, we consider the following generalization, also called *Lin's index*, which is quite established in the literature (Lin 1976; Wasserman and Faust 1994; Boldi

²<https://networkit.github.io/>.

Table 1. Notations Used Throughout the Article

Symbol	Definition
Graphs	
$G = (V, E)$	Graph with node/vertex set V and edge/arc set E
n, m	$ V , E $
$\mathcal{G} = (\mathcal{V}, \mathcal{E}, w)$	Weighted directed acyclic graph of strongly connected components (see Section 8.4)
$\deg(v)$	Degree of a node in an undirected graph
$\text{outdeg}(v)$	Out-degree of a node in a directed graph
$d(v, w)$	Number of edges in a shortest path from v to w
Reachability set function	
$R(v)$	Set of nodes reachable from v (by definition, $v \in R(v)$)
$r(v)$	$ R(v) $
$\alpha(v)$	Lower bound on $r(v)$, that is, $\alpha(v) \leq r(v)$ (see Section 8.4)
$\omega(v)$	Upper bound on $r(v)$, that is, $r(v) \leq \omega(v)$ (see Section 8.4)
Neighborhood functions	
$\Gamma_d(v)$	Set of nodes at distance d from v : $\{w \in V : d(v, w) = d\}$
$\Gamma(v)$	Set of neighbors of v , that is, $\Gamma_1(v)$
$\gamma_d(v)$	Number of nodes at distance d from v , that is, $ \Gamma_d(v) $
$\tilde{\gamma}_d(v)$	Upper bound on $\gamma_d(v)$ computed using the neighborhood-based lower bound (see Section 5)
$\tilde{\gamma}_{d+1}(v)$	Upper bound on $\gamma_{d+1}(v)$, defined as $\sum_{u \in \Gamma_d(v)} \deg(u) - 1$ if the graph is undirected, $\sum_{u \in \Gamma_d(v)} \text{outdeg}(u)$ otherwise
$N_d(v)$	Set of nodes at distance <i>at most</i> d from v , that is, $\{w \in V : d(v, w) \leq d\}$
$n_d(v)$	Number of nodes at distance <i>at most</i> d from v , that is, $ N_d(v) $
Closeness functions	
$c(v)$	Closeness of node v , that is, $\frac{(r(v)-1)^2}{(n-1) \sum_{w \in R(v)} d(v, w)}$
Distance sum functions	
$S(v)$	Total distance of node v , that is $\sum_{w \in R(v)} d(v, w)$
$S^{\text{NB}}(v, r)$	Lower bound on $S(v)$ if $r(v) = r$, used in the computeBoundsNB function (see Proposition 5.1)
$S_d^{\text{CUT}}(v, r)$	Lower bound on $S(v)$ if $r(v) = r$, used in the updateBoundsBFSCut function (see Lemma 6.1)
$S_s^{\text{LB}}(v, r)$	Lower bound on $S(v)$ if $r(v) = r$, used in the updateBoundsLB function (see Equation (4), (5))
Farness functions	
$f(v)$	Farness of node v , that is, $\frac{(n-1)S(v)}{(R(v) -1)^2}$
$L(v, r)$	Generic lower bound on $f(v)$, if $r(v) = r$
$L^{\text{NB}}(v, r)$	Lower bound on $f(v)$, if $r(v) = r$, defined as $(n-1) \frac{S^{\text{NB}}(v, r)}{(r-1)^2}$
$L_d^{\text{CUT}}(v, r)$	Lower bound on $f(v)$, if $r(v) = r$, defined as $(n-1) \frac{S_d^{\text{CUT}}(v, r)}{(r-1)^2}$
$L_s^{\text{LB}}(v, r)$	Lower bound on $f(v)$, if $r(v) = r$, defined as $(n-1) \frac{S_s^{\text{LB}}(v, r)}{(r-1)^2}$

and Vigna 2013, 2014; Olsen et al. 2014):

$$f(v) = \frac{\sum_{w \in R(v)} d(v, w)}{r(v) - 1} \cdot \frac{n - 1}{r(v) - 1} \quad c(v) = \frac{1}{f(v)} \quad (1)$$

If a vertex v has (out)degree 0, the previous fraction becomes $\frac{0}{0}$: in this case, the closeness of v is set to 0.

With this definition, if a vertex v has only one neighbor w at distance 1, then its closeness is $\frac{1}{n-1}$, which is much smaller than what we would obtain with the naive generalization, and it better matches our intuition.

Another possibility is to consider a slightly different definition:

$$c(v) = \sum_{w \in V} f(d(v, w)),$$

for some decreasing function f .³ One of the most common choices of f is $f(d) = \frac{1}{d}$: this way, we obtain the harmonic centrality (Marchiori and Latora 2000).

In this article, we focus on Lin's index, because it is quite established in the literature, because the previously best exact top- k closeness centrality algorithm uses this definition (Olsen et al. 2014), and because, when restricted to the connected case, this definition coincides with closeness centrality (from now on, in a disconnected context, we use closeness centrality to indicate Lin's index). However, all our algorithms can be easily adapted to any centrality measure of the form $c(v) = \sum_{w \in V} f(d(v, w))$: indeed, in Section 9, we show that our algorithm performs very well also with harmonic centrality.

3 COMPLEXITY OF COMPUTING THE MOST CENTRAL VERTEX

In this section, we show that, even in the computation of the most central vertex, the *textbook* algorithm is almost optimal in the worst case, assuming the Orthogonal Vector conjecture (Williams 2005; Abboud et al. 2016), or the well-known SETH (Impagliazzo et al. 2001).

We recall that the Orthogonal Vector conjecture states that, given N vectors in $\{0, 1\}^d$ (where $d = O(\log^k N)$ for some k) it is impossible to decide whether there are two orthogonal vectors in $O(N^{2-\epsilon})$ time, for any $\epsilon > 0$ not depending on k .

SETH states that the k -SATISFIABILITY problem cannot be solved in time $O((2-\epsilon)^N)$, where N is the number of variables and ϵ is a positive constant not depending on k . It is well known that the Orthogonal Vector conjecture is implied by SETH (Williams 2005; Borassi et al. 2015a; Abboud et al. 2016).

Our reduction is summarized by the following theorem, which shows the hardness of computing the most central vertex on directed, non-strongly connected graphs.

THEOREM 3.1. *On directed sparse graphs (i.e., graphs for which $m = O(n)$), an algorithm computing the most closeness central vertex in time $O(n^{2-\epsilon})$ for some $\epsilon > 0$ would falsify the Orthogonal Vector conjecture.*

It is worth mentioning that this result still holds if we restrict our analysis to graphs with small diameter. Indeed, the diameter of the graph obtained from the reduction is 9 (see Appendix A).

Theorem 3.1 directly implies the following corollary.

COROLLARY 3.2. *On directed graphs, an algorithm computing the most closeness central vertex in time $O(m^{2-\epsilon})$ for some $\epsilon > 0$ would falsify the Orthogonal Vector conjecture.*

PROOF. When applied to sparse graphs, such an algorithm would run in $O(m^{2-\epsilon}) = O(n^{2-\epsilon})$, against Theorem 3.1. \square

Since the Orthogonal Vector conjecture is implied by SETH, our results still hold if we replace the Orthogonal Vector conjecture with SETH in the statement of Theorem 3.1, Corollary 3.2.

The proof of Theorem 3.1 is given in Appendix A.

³Usually, it is also assumed without loss of generality that $f(+\infty) = 0$, that is, we consider only reachable vertices: if this is not the case, it is enough to use a new function defined by $g(d) = f(d) - f(+\infty)$.

4 OVERVIEW OF THE ALGORITHM

In this section, we describe our new approach for computing the k nodes with maximum closeness (equivalently, the k nodes with minimum farness, where the farness $f(v)$ of a vertex v is $\frac{1}{c(v)} = \frac{(n-1) \sum_{w \in R(v)} d(v, w)}{(r(v)-1)^2}$, as in Table 1). If we have more than one node with the same score, we output all nodes having a centrality bigger than or equal to the centrality of the k th node.

In the previous section, we have shown that the trivial algorithm cannot be improved in the worst case: here, we describe an algorithm that is much more efficient when tested on real-world graphs. The basic idea is to keep track of a lower bound on the farness of each node, and to skip the analysis of a vertex v if this lower bound implies that v is not in the top k .

More formally, let us assume that we know the farness of some vertices v_1, \dots, v_l , and a lower bound $L(w)$ on the farness of any other vertex w . Furthermore, assume that there are k vertices among v_1, \dots, v_l verifying $f(v_i) < L(w) \forall w \in V - \{v_1, \dots, v_l\}$. This means that, for any of these vertices v_i , it holds that $f(w) \geq L(w) > f(v_i) \forall w \in V - \{v_1, \dots, v_l\}$. Then, we can safely skip the exact computation of $f(w)$ for all remaining nodes w , because the k vertices with smallest farness are among v_1, \dots, v_l .

This idea is implemented in Algorithm 1: we use a list `Top` containing all ‘analyzed’ vertices v_1, \dots, v_l in increasing order of farness, and a priority queue `Q` containing all vertices “not analyzed, yet,” in increasing order of lower bound L (this way, the head of `Q` always has the smallest value of L among all vertices in `Q`). At the beginning, using the function `computeBounds()`, we compute a first bound $L(v)$ for each vertex v , and we fill the queue `Q` according to this bound. Then, at each step, we extract the first element v of `Q`: if $L(v)$ is higher than the k th smallest farness computed until now (that is, the farness of the k th vertex in variable `Top`), we can safely stop, because for each $x \in Q$, $f(x) \geq L(x) \geq L(v) > f(\text{Top}[k])$, and x is not in the top k . Otherwise, we run the function `updateBounds(v)`, which performs a BFS from v , returns the farness of v , and improves the bounds L of all other vertices. Finally, we insert v into `Top` in the right position, and we update `Q` if the lower bounds have changed.

The crucial point of the algorithm is the definition of the lower bounds, that is, the definition of the functions `computeBounds` and `updateBounds`. We propose two alternative strategies for each of these two functions (see Figure 1): in both cases, one strategy is conservative, that is, it tries to perform as few operations as possible, while the other strategy is aggressive, namely, it needs many operations, but at the same time it improves many lower bounds.

Let us analyze the possible choices of the function `computeBounds`. The conservative strategy `computeBoundsDeg` needs time $O(n)$: it simply sets $L(v) = 0$ for each v , and it fills `Q` by inserting

ALGORITHM 1: Pseudocode of our algorithm for top- k closeness centralities.

```

Input: A graph  $G = (V, E)$ 
Output: Top  $k$  nodes with highest closeness
1 global L, Q  $\leftarrow$  computeBounds( $G$ );
2 global Top  $\leftarrow$  [];
3 global Farn;
4 for  $v \in V$  do Farn[ $v$ ] =  $+\infty$ ;
5 while Q is not empty do
6    $v \leftarrow$  Q.extractMin();
7   if |Top|  $\geq k$  and L[ $v$ ] > Farn[Top[ $k$ ]] then return Top[1 :  $k$ ];
8   Farn[ $v$ ]  $\leftarrow$  updateBounds( $v$ , L); // This function might also modify L
9   add  $v$  to Top, and sort Top according to Farn;
10  update Q according to the new bounds;
```

ALGORITHM 1: Pseudocode of our algorithm for top- k closeness centralities.

Input : A graph $G = (V, E)$
Output: Top k nodes with highest closeness

```

1 global L, Q ← computeBounds(G);
2 global Top ← [];
3 global Farn;
4 for  $v \in V$  do Farn[v] =  $+\infty$ ;
5 while Q is not empty do
6    $v \leftarrow Q.extractMin()$ ;
7   if  $|Top| \geq k$  and  $L[v] > Farn[Top[k]]$  then return Top[1 : k];
8   Farn[v] ← updateBounds(v, L); // This function might also modify L
9   add v to Top, and sort Top according to Farn;
10  update Q according to the new bounds;
```

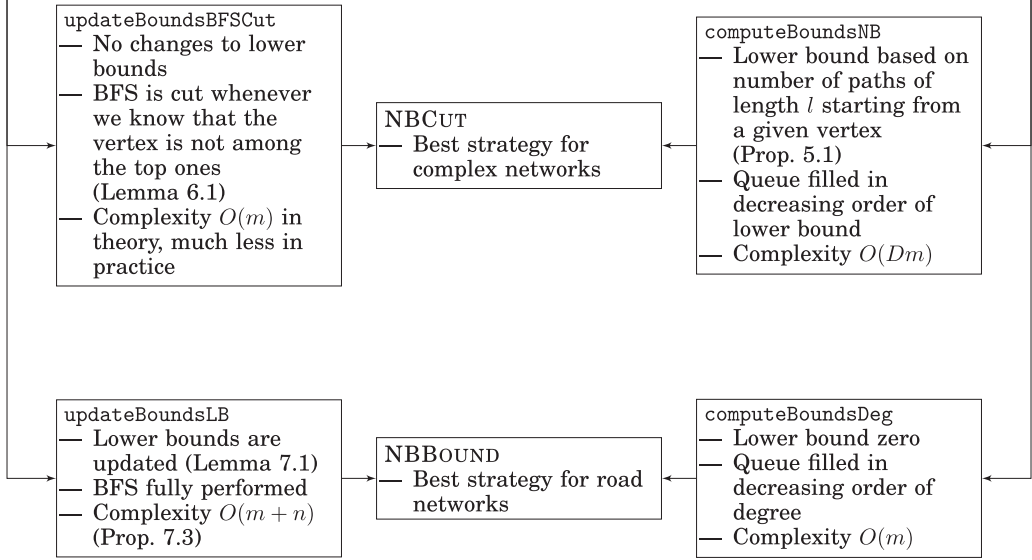


Fig. 1. The two alternative strategies for each of the two functions `computeBounds` and `updateBounds`: in both cases, the first strategy is conservative (that is, it tries to perform as few operations as possible), while the other strategy is aggressive (namely, it needs many operations, but at the same time it improves many lower bounds). In the middle of the figure, we show the two combinations of these four strategies, which perform better on real-word networks.

nodes in decreasing order of degree (the idea is that vertices with high degree have small farness, and they should be analyzed as early as possible, so that the values in `Top` are correct as soon as possible). Note that the vertices can be sorted in time $O(n)$ using counting sort.

The aggressive strategy `computeBoundsNB` needs time $O(mD)$, where D is the diameter of the graph: it computes the neighborhood-based lower bound $L^{NB}(v)$ for each vertex v (we will explain shortly afterwards how it works), it sets $L(v) = L^{NB}(v)$, and it fills `Q` by adding vertices in decreasing order of L . The idea behind the neighborhood-based lower bound is to count the number of paths of length l starting from a given vertex v , which is also an upper bound U_l on the number of vertices at distance l from v . From U_l , it is possible to define a lower bound on $\sum_{x \in V} d(v, x)$ by “summing U_l times the distance l ,” until we have summed n distances: this bound yields the desired lower bound on the farness of v . The detailed explanation of this function is provided in Section 5.

For the function `updateBounds(v)`, the conservative strategy `updateBoundsBFSCut(v)` does not improve L , and it cuts the BFS as soon as it is sure that the farness of v is larger than the k th smallest farness found until now, that is, $Farn[Top[k]]$. If the BFS is cut, the function returns $+\infty$, otherwise, at the end of the BFS we have computed the farness of v , and we can return it. The running time

of this procedure is $O(m)$ in the worst case, but it can be much better in practice. It remains to define how the procedure can be sure that the farness of v is at least x : to this purpose, during the BFS, we update a lower bound on the farness of v . The idea behind this bound is that, if we have already visited all nodes up to distance d , we can lower bound the farness of v by setting distance $d + 1$ to a number of vertices equal to the number of edges “leaving” level d , and distance $d + 2$ to all the remaining vertices. The details of this procedure are provided in Section 6.

The aggressive strategy `updateBoundsLB`(v) performs a complete BFS from v , and it bounds the farness of each node w using the level-based lower bound. The running time is $O(m)$ for the BFS, and $O(n)$ to compute the bounds. The idea behind the level-based lower bound is that, for each two nodes w and x , $d(w, x) \geq |d(v, w) - d(v, x)|$, and consequently $\sum_{x \in V} d(w, x) \geq \sum_{x \in V} |d(v, w) - d(v, x)|$. The latter sum can be computed in time $O(n)$ for each w , because it depends only on the level d of w in the BFS tree, and because it is possible to compute in $O(1)$ the sum for a vertex at level $d + 1$, if we know the sum for a vertex at level d . The details are provided in Section 7.

Finally, in order to transform these lower bounds on $\sum_{x \in V} d(v, x)$ into bounds on $f(v)$, we need to know the number of vertices reachable from a given vertex v . In Sections 5–7, we assume that these values are known: this assumption is true in undirected graphs, where we can compute the number of reachable vertices in linear time at the beginning of the algorithm, and in strongly connected directed graphs, where the number of reachable vertices is n . The only remaining case is when the graph is directed and not strongly connected: in this case, we need some additional machinery, which are presented in Section 8.

5 NEIGHBORHOOD-BASED LOWER BOUND

In this section, we propose a lower bound $S^{\text{NB}}(v, r(v))$ on the total sum $S(v) = \sum_{w \in R(v)} d(v, w)$. For simplicity, we assume for now that the graph is undirected or strongly-connected (in Section 8 we will discuss how these results can be extended to general graphs). Notice that if we know the number $r(v)$ of vertices reachable from v , a bound on the total sum directly translates into a lower bound on the farness of v , simply multiplying by $(n - 1)/(r(v) - 1)^2$.

The basic idea is to find an upper bound $\tilde{y}_i(v)$ on the number of nodes $y_i(v)$ at distance i from v . Then, intuitively, if we assume that the number of nodes at distance i is greater than its actual value and “stop counting” when we have $r(v)$ nodes, we get something that is smaller than the actual total distance. This is because we are assuming that the distances of some nodes are smaller than their actual values. This argument is formalized in Proposition 5.1.

PROPOSITION 5.1. *If $\tilde{y}_i(v)$ is an upper bound on $y_i(v)$, for $i = 0, \dots, \text{diam}(G)$ and $\text{ecc}(v) := \max_{w \in r(v)} d(v, w)$, then $S^{\text{NB}}(v, r(v)) := \sum_{k=1}^{\text{ecc}(v)} k \cdot \min\{\tilde{y}_k(v), \max\{r(v) - \sum_{i=0}^{k-1} \tilde{y}_i(v), 0\}\}$ is a lower bound on $S(v)$.*

PROOF. First, we notice that $S(v) = \sum_{k=0}^{\text{ecc}(v)} k \cdot \gamma_k(v)$ and $r(v) = \sum_{k=0}^{\text{ecc}(v)} \gamma_k(v)$.

Let us assume that $\tilde{y}_0(v) < r(v)$. In fact, if $\tilde{y}_0(v) \geq r(v)$, the statement is trivially satisfied. Then, there must be a number $\text{ecc}' > 0$ such that for $k < \text{ecc}'$ the quantity $\min\{\tilde{y}_k(v), \max\{r(v) - \sum_{i=0}^{k-1} \tilde{y}_i(v), 0\}\}$ is equal to $\tilde{y}_k(v)$, for $k = \text{ecc}'$, the quantity is equal to $\alpha := r(v) - \sum_{k=0}^{\text{ecc}'-1} \tilde{y}_k(v) > 0$ and, for $k > \text{ecc}'$, it is equal to 0. Therefore we can write $S^{\text{NB}}(v, r(v))$ as $\sum_{k=1}^{\text{ecc}'-1} k \cdot \tilde{y}_k(v) + \text{ecc}' \cdot \alpha$.

We show that $\text{ecc}' \leq \text{ecc}(v)$. In fact, we know that $\sum_{k=0}^{\text{ecc}'-1} \tilde{y}_k(v) < r(v) = \sum_{k=0}^{\text{ecc}(v)} \gamma_k(v) \leq \sum_{k=0}^{\text{ecc}(v)} \tilde{y}_k(v)$. Therefore $\text{ecc}' - 1 < \text{ecc}(v)$, which implies $\text{ecc}' \leq \text{ecc}(v)$.

For each i , we can write $\tilde{y}_i(v) = y_i(v) + \epsilon_i$, $\epsilon_i \geq 0$. Therefore, we can write $\sum_{k=0}^{\text{ecc}'-1} \epsilon_i + \alpha = r(v) - \sum_{k=0}^{\text{ecc}'-1} \gamma_k(v) = \sum_{k=\text{ecc}'}^{\text{ecc}(v)} \gamma_k(v)$. Then, $S^{\text{NB}}(v, r(v)) = \sum_{k=0}^{\text{ecc}'-1} k \cdot \gamma_k(v) + \sum_{k=0}^{\text{ecc}'-1} k \cdot \epsilon_i + \text{ecc}' \cdot \alpha \leq \sum_{k=0}^{\text{ecc}'-1} k \cdot \gamma_k(v) + \text{ecc}'(\alpha + \sum_{k=0}^{\text{ecc}'-1} \epsilon_i) = \sum_{k=0}^{\text{ecc}'-1} k \cdot \gamma_k(v) + \text{ecc}'(\sum_{k=\text{ecc}'}^{\text{ecc}(v)} \gamma_k(v)) \leq \sum_{k=0}^{\text{ecc}(v)} k \cdot \gamma_k(v) = S(v)$. \square

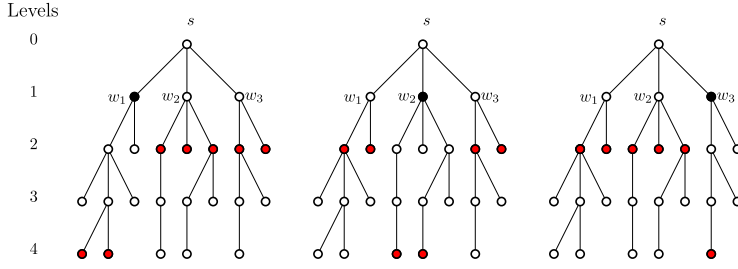


Fig. 2. Relation between nodes at distance 4 for s and the neighbors of s . The red nodes represent the nodes at distance 3 for w_1 (left), for w_2 (center), and for w_3 (right).

In the following paragraphs, we propose upper bounds $\tilde{\gamma}_i(v)$ for trees, undirected graphs and directed strongly-connected graphs. In case of trees, the bound $\tilde{\gamma}_i(v)$ is actually equal to $\gamma_i(v)$, which means that the algorithm can be used to compute closeness of all nodes in a tree exactly.

Computing closeness on trees. Let us consider a node s for which we want to compute the total distance $S(s)$ (notice that in a tree $c(s) = (n - 1)/S(s)$). The number of nodes at distance 1 in the BFS tree from s is clearly the degree of s . What about distance 2? Since there are no cycles, all the neighbors of the nodes in $\Gamma_1(s)$ are nodes at distance 2 from s , with the only exception of s itself. Therefore, naming $\Gamma_k(s)$ the set of nodes at distance k from s and $\gamma_k(s)$ the number of these nodes, we can write $\gamma_2(s) = \sum_{w \in \Gamma_1(s)} \gamma_1(w) - \deg(s)$. In general, we can always relate the number of nodes at each distance k of s to the number of nodes at distance $k - 1$ in the BFS trees of the neighbors of s . Let us now consider $\gamma_k(s)$, for $k > 2$.

Figure 2 shows an example where s has three neighbors w_1 , w_2 and w_3 . Suppose we want to compute $\Gamma_4(s)$ using information from w_1 , w_2 and w_3 . Clearly, $\Gamma_4(s) \subset \Gamma_3(w_1) \cup \Gamma_3(w_2) \cup \Gamma_3(w_3)$; however, there are also other nodes in the union that are not in $\Gamma_4(s)$. Furthermore, the nodes in $\Gamma_3(w_1)$ (red nodes in the leftmost tree) are of two types: nodes in $\Gamma_4(s)$ (the ones in the subtree of w_1) and nodes in $\Gamma_2(s)$ (the ones in the subtrees of w_2 and w_3). An analogous behavior can be observed for w_2 and w_3 (central and rightmost trees). If we simply sum all the nodes in $\gamma_3(w_1)$, $\gamma_3(w_2)$ and $\gamma_3(w_3)$, we would be counting each node at level 2 twice, i.e., once for each node in $\Gamma_1(s)$ minus one. Hence, for each $k > 2$, we can write

$$\gamma_k(s) = \sum_{w \in \Gamma_1(s)} \gamma_{k-1}(w) - \gamma_{k-2}(s) \cdot (\deg(s) - 1). \quad (2)$$

From this observation, we define a new method to compute the total distance of all nodes, described in Algorithm 2. Instead of computing the BFS tree of each node one by one, at each step we compute the number $\gamma_k(v)$ of nodes at level k for *all* nodes v . First (Lines 2–4), we compute $\gamma_1(v)$ for each node (and add that to $S(v)$). Then (Lines 6–19), we consider all the other levels k one by one. For each k , we use $\gamma_{k-1}(w)$ of the neighbors w of v and $\gamma_{k-2}(v)$ to compute $\gamma_k(v)$ (Lines 9 and 11). If, for some k , $\gamma_k(v) = 0$, all the nodes have been added to $S(v)$. Therefore, we can stop the algorithm when $\gamma_k(v) = 0 \ \forall v \in V$.

PROPOSITION 5.2. *Algorithm 2 requires $O(D \cdot n)$ operations to compute the closeness centrality of all nodes in a tree T .*

PROOF. The for loop in Lines 2–4 of Algorithm 2 clearly takes $O(n)$ time. For each level of the while loop of Lines 6–19, each node scans its neighbors in Line 9 or Line 11. In total, this leads to $O(n)$ operations per level since $m = O(n)$. Since the maximum number of levels that a node can have is equal to the diameter of the tree, the algorithm requires $O(D \cdot n)$ operations. \square

ALGORITHM 2: Closeness centrality in trees

Input: A tree $T = (V, E)$
Output: Closeness centralities $c(v)$ of each node $v \in V$

```

1   $k \leftarrow 2$ ;
2  foreach  $s \in V$  do
3     $\gamma_{k-1}(s) \leftarrow \deg(s)$ ;
4     $S(s) \leftarrow \deg(s)$ ;
5   $nFinished \leftarrow 0$ ;
6  while  $nFinished < n$  do
7    foreach  $s \in V$  do
8      if  $k = 2$  then
9         $\gamma_k(s) \leftarrow \sum_{w \in \Gamma(s)} \gamma_{k-1}(w) - \deg(s)$ ;
10     else
11        $\gamma_k(s) \leftarrow \sum_{w \in \Gamma(s)} \gamma_{k-1}(w) - \gamma_{k-2}(s)(\deg(s) - 1)$ ;
12     foreach  $s \in V$  do
13        $\gamma_{k-2}(s) \leftarrow \gamma_{k-1}(s)$ ;
14        $\gamma_{k-1}(s) \leftarrow \gamma_k(s)$ ;
15       if  $\gamma_{k-1}(s) > 0$  then
16          $S(s) \leftarrow S(s) + k \cdot \gamma_{k-1}(s)$ ;
17       else
18          $nFinished \leftarrow nFinished + 1$ ;
19    $k \leftarrow k + 1$ ;
20 foreach  $s \in V$  do
21    $c(s) \leftarrow (n - 1) / S(s)$ ;
22 return  $c$ 

```

Note that closeness centrality on trees could even be computed in time $O(n)$ in a different manner (Brandes and Fleischer 2005). We choose to include Algorithm 2 here nonetheless since it paves the way for an algorithm computing a lower bound in general undirected graphs, described next.

Lower bound for undirected graphs. For general undirected graphs, Equation (2) is not true anymore – but a related upper bound $\tilde{\gamma}_k(\cdot)$ on $\gamma_k(\cdot)$ is still useful. Let $\tilde{\gamma}_k(s)$ be defined recursively as in Equation (2): in a tree, $\tilde{\gamma}_k(s) = \gamma_k(s)$, while in this case we prove that $\tilde{\gamma}_k(s)$ is an upper bound on $\gamma_k(s)$. Indeed, there could be nodes x for which there are multiple paths between s and x and that are therefore contained in the subtrees of more than one neighbor of s . This means that we would count x multiple times when considering $\tilde{\gamma}_k(s)$, overestimating the number of nodes at distance k . However, we know for sure that at level k there cannot be *more nodes* than in Equation (2). If, for each node v , we assume that the number $\tilde{\gamma}_k(v)$ of nodes at distance k is that of Equation (2), we can apply Proposition 5.1 and get a lower bound $S^{NB}(v, r(v))$ on the total sum for undirected graphs. The procedure is described in Algorithm 3. The computation of $S^{NB}(v, r(v))$ works basically like Algorithm 2, with the difference that here we keep track of the number of the nodes found in all the levels up to k ($nVisited$) and stop the computation when $nVisited$ becomes equal to $r(v)$ (if it becomes larger, in the last level we consider only $r(v) - nVisited$ nodes, as in Proposition 5.1 (Lines 22–25)).

PROPOSITION 5.3. *For an undirected graph G , computing the lower bound $S^{NB}(v, r(v))$ described in Algorithm 3 takes $O(D \cdot m)$ time.*

PROOF. Like in Algorithm 2, the number of operations performed by Algorithm 3 at each level of the while loop is $O(m)$. At each level i , all the nodes at distance i are accounted for (possibly

multiple times) in Lines 11 and 13. Therefore, at each level, the variable $nVisited$ is always greater than or equal to the the number of nodes v at distance at most i . Since this distance is at most D for all nodes v , the maximum number of levels scanned in the while loop cannot be larger than D , therefore the total complexity is $O(D \cdot m)$. \square

ALGORITHM 3: Neighborhood-based lower bound for undirected graphs

Input: A graph $G = (V, E)$
Output: Lower bounds $L^{NB}(v, r(v))$ of each node $v \in V$

```

1  $k \leftarrow 2$ ;
2 foreach  $s \in V$  do
3    $\gamma_{k-1}(s) \leftarrow \deg(s)$ ;
4    $\tilde{S}^{(un)}(s) \leftarrow \deg(s)$ ;
5    $nVisited[s] \leftarrow \deg(s) + 1$ ;
6    $finished[s] \leftarrow false$ ;
7  $nFinished \leftarrow 0$ ;
8 while  $nFinished < n$  do
9   foreach  $s \in V$  do
10    if  $k = 2$  then
11       $\gamma_k(s) \leftarrow \sum_{w \in \Gamma(s)} \gamma_{k-1}(w) - \deg(s)$ ;
12    else
13       $\gamma_k(s) \leftarrow \sum_{w \in \Gamma(s)} \gamma_{k-1}(w) - \gamma_{k-2}(s)(\deg(s) - 1)$ ;
14  foreach  $s \in V$  do
15    if  $finished[s]$  then
16      continue;
17     $\gamma_{k-2}(s) \leftarrow \gamma_{k-1}(s)$ ;
18     $\gamma_{k-1}(s) \leftarrow \gamma_k(s)$ ;
19     $nVisited[s] \leftarrow nVisited[s] + \gamma_{k-1}(s)$ ;
20    if  $nVisited[s] < r(v)$  then
21       $\tilde{S}^{(un)}(s) \leftarrow \tilde{S}^{(un)}(s) + k \cdot \gamma_{k-1}(s)$ ;
22    else
23       $\tilde{S}^{(un)}(s) \leftarrow \tilde{S}^{(un)}(s) + k(r(v) - (nVisited[s] - \gamma_{k-1}(s)))$ ;
24       $nFinished \leftarrow nFinished + 1$ ;
25       $finished[s] \leftarrow true$ ;
26   $k \leftarrow k + 1$ ;
27 foreach  $v \in v$  do
28    $L^{NB}(v, r(v)) \leftarrow \frac{(n-1)\tilde{S}^{(un)}}{(r(v)-1)^2}$ ;
29 return  $L^{NB}(\cdot, r(\cdot))$ 

```

Lower bound on directed graphs. In directed graphs, we can simply consider the out-neighbors, without subtracting the number of nodes discovered in the subtrees of the other neighbors in Equation (2). The lower bound (which we still refer to as $S^{NB}(v, r(v))$) is obtained by replacing Equation (2) with the following in Lines 11 and 13 of Algorithm 3:

$$\tilde{\gamma}_k(s) = \sum_{w \in \Gamma(s)} \tilde{\gamma}_{k-1}(w) \quad (3)$$

6 THE UPDATEBOUNDSBFSCUT FUNCTION

The `updateBoundsBFSCut` function is based on a simple idea: if the k th biggest farness found until now is x , and if we are performing a BFS from vertex v to compute its farness $f(v)$, we can stop as soon as we can guarantee that $f(v) \geq x$.

Informally, assume that we have already visited all nodes up to distance d : we can lower bound $S(v) = \sum_{w \in V} d(v, w)$ by setting distance $d + 1$ to a number of vertices equal to the number of edges “leaving” level d , and distance $d + 2$ to all the remaining reachable vertices. Then, this bound yields a lower bound on the farness of v . As soon as this lower bound is bigger than x , the `updateBoundsBFSCut` function may stop; if this condition never occurs, at the end of the BFS we have exactly computed the farness of x .

More formally, the following lemma defines a lower bound $S_d^{\text{CUT}}(v, r(v))$ on $S(v)$, which is computable after we have performed a BFS from v up to level d , assuming we know the number $r(v)$ of vertices reachable from v (this assumption is lifted in Section 8).

LEMMA 6.1. *Given a graph $G = (V, E)$, a vertex $v \in V$, and an integer $d \geq 0$, let $N_d(v)$ be the set of vertices at distance at most d from v , $n_d(v) = |N_d(v)|$, and let $\tilde{\gamma}_{d+1}(v)$ be an upper bound on the number of vertices at distance $d + 1$ from v (see Table 1). Then,*

$$S(v) \geq S_d^{\text{CUT}}(v, r(v)) := \sum_{w \in N_d(v)} d(v, w) - \tilde{\gamma}_{d+1}(v) + (d + 2)(r(v) - n_d(v)).$$

PROOF. The sum of all the distances from v is lower bounded by setting the correct distance to all vertices at distance at most d from v , by setting distance $d + 1$ to all vertices at distance $d + 1$ (there are $\gamma_{d+1}(v)$ such vertices), and by setting distance $d + 2$ to all other vertices (there are $r(v) - n_{d+1}(v)$ such vertices, where $r(v)$ is the number of vertices reachable from v and $n_{d+1}(v)$ is the number of vertices at distance at most $d + 1$). More formally, $f(v) \geq \sum_{w \in N_d(v)} d(v, w) + (d + 1)\gamma_{d+1}(v) + (d + 2)(r(v) - n_{d+1}(v))$.

Since $n_{d+1}(v) = \gamma_{d+1}(v) + n_d(v)$, we obtain that $f(v) \geq \sum_{w \in N_d(v)} d(v, w) - \gamma_{d+1}(v) + (d + 2)(r(v) - n_d(v))$. We conclude because, by assumption, $\tilde{\gamma}_{d+1}(v)$ is an upper bound on $\gamma_{d+1}(v)$. \square

COROLLARY 6.2. *For each vertex v and for each $d \geq 0$,*

$$f(v) \geq L_d^{\text{CUT}}(v, r(v)) := \frac{(n - 1)S_d^{\text{CUT}}(v, r(v))}{(r(v) - 1)^2}.$$

It remains to define the upper bound $\tilde{\gamma}_{d+1}(v)$: in the directed case, this bound is simply the sum of the out-degrees of vertices at distance d from v . In the undirected case, since at least an edge from each vertex $v \in \Gamma_d(v)$ is directed towards $\Gamma_{d-1}(v)$, we may define $\tilde{\gamma}_{d+1}(v) = \sum_{w \in \Gamma_d(v)} \deg(w) - 1$ (the only exception is $d = 0$: in this case, $\tilde{\gamma}_1(v) = \gamma_1(v) = \deg(v)$).

Remark 6.3. When we are processing vertices at level d , if we process an edge (x, y) where y is already in the BFS tree, we can decrease $\tilde{\gamma}_{d+1}(v)$ by one, obtaining a better bound.

Assuming we know $r(v)$, all quantities necessary to compute $L_d^{\text{CUT}}(v, r(v))$ are available as soon as all vertices in $N_d(v)$ are visited by a BFS. This function performs a BFS starting from v , continuously updating the upper bound $L_d^{\text{CUT}}(v, r(v)) \leq f(v)$ (the update is done whenever all nodes in $\Gamma_d(v)$ have been reached, or Remark 6.3 can be used). As soon as $L_d^{\text{CUT}}(v, r(v)) \geq x$, we know that $f(v) \geq L_d^{\text{CUT}}(v, r(v)) \geq x$, and we return $+\infty$, meaning that v is not one of the k most central nodes.

Algorithm 4 is the pseudocode of the function `updateBoundsBFSCut` when implemented for directed graphs, assuming we know the number $r(v)$ of vertices reachable from each v (for example, if the graph is strongly connected). This code can be easily adapted to all the other cases.

7 THE UPDATEBOUNDSLB FUNCTION

Differently from `updateBoundsBFSCut` function, `updateBoundsLB` computes a complete BFS traversal, but uses information acquired during the traversal to update the bounds on the other nodes.

ALGORITHM 4: The `updateBoundsBFSCut(v)` function in the case of directed graphs, if $r(v)$ is known for each v .

```

1   $x \leftarrow \text{Farn}(\text{Top}[k]);$  // Farn and Top are global variables, as in Algorithm 1.
2  Create queue  $Q$ ;
3   $Q.\text{enqueue}(v);$ 
4  Mark  $v$  as visited;
5   $d \leftarrow 0; S \leftarrow 0; \tilde{\gamma} \leftarrow \text{outdeg}(v); nd \leftarrow 1;$ 
6  while  $Q$  is not empty do
7       $u \leftarrow Q.\text{dequeue}();$ 
8      if  $d(v, u) > d$  then
9           $d \leftarrow d + 1;$ 
10          $L_d^{\text{CUT}}(v, r(v)) \leftarrow \frac{(n-1)(S-\tilde{\gamma}+(d+2)(r(v)-nd))}{(r(v)-1)^2};$ 
11         if  $L_d^{\text{CUT}}(v, r(v)) \geq x$  then return  $+\infty;$ 
12          $\tilde{\gamma} \leftarrow 0$ 
13     for  $w$  in adjacency list of  $u$  do
14         if  $w$  is not visited then
15              $S \leftarrow S + d(v, w);$ 
16              $\tilde{\gamma} \leftarrow \tilde{\gamma} + \text{outdeg}(w);$ 
17              $nd \leftarrow nd + 1;$ 
18              $Q.\text{enqueue}(w);$ 
19             Mark  $w$  as visited
20         else
21             // we use Remark 6.3
22              $L_d^{\text{CUT}}(v, r(v)) \leftarrow L_d^{\text{CUT}}(v, r(v)) + \frac{(n-1)}{(r(v)-1)^2};$ 
23             if  $L_d^{\text{CUT}}(v, r(v)) \geq x$  then return  $+\infty;$ 
24 return  $\frac{S(n-1)}{(r(v)-1)^2};$ 

```

Let us first consider an undirected graph G and let s be the source node from which we are computing the BFS. We can see the distances $d(s, v)$ between s and all the nodes v reachable from s as *levels*: node v is at level i if and only if the distance between s and v is i , and we write $v \in \Gamma_i(s)$ (or simply $v \in \Gamma_i$ if s is clear from the context). Let i and j be two levels, $i \leq j$. Then, the distance between any two nodes v at level i and w at level j must be at least $j - i$. Indeed, if $d(v, w)$ was smaller than $j - i$, w would be at level $i + d(v, w) < j$, which contradicts our assumption. It follows directly that $\sum_{w \in V} |d(s, w) - d(s, v)|$ is a lower bound on $S(v)$, for all $v \in R(s)$:

LEMMA 7.1. $\sum_{w \in R(s)} |d(s, w) - d(s, v)| \leq S(v) \quad \forall v \in R(s).$

To improve the approximation, we notice that the number of nodes at distance 1 from v is exactly the degree of v . Therefore, all the other nodes w such that $|d(s, v) - d(s, w)| \leq 1$ must be at least at distance 2 (with the only exception of v itself, whose distance is of course 0). This way we can define the following lower bound on $S(v)$:

$$\begin{aligned}
 & 2(\#\{w \in R(s) : |d(s, w) - d(s, v)| \leq 1\} - \deg(v) - 1) \\
 & + \deg(v) + \sum_{\substack{w \in R(s) \\ |d(s, w) - d(s, v)| > 1}} |d(s, w) - d(s, v)|,
 \end{aligned}$$

that is,

$$2 \cdot \sum_{|j-d(s,v)| \leq 1} \gamma_j + \sum_{|j-d(s,v)| > 1} \gamma_j \cdot |j - d(s,v)| - \deg(v) - 2, \quad (4)$$

where $\gamma_j = |\Gamma_j|$.

Multiplying the bound of Equation (4) by $\frac{(n-1)}{(r(v)-1)^2}$, we obtain a lower bound on the fairness $f(v)$ of node v , named $L_s^{\text{LB}}(v, r(v))$. A straightforward way to compute $L_s^{\text{LB}}(v, r(v))$ would be to first run the BFS from s and then, for each node v , to consider the level difference between v and all the other nodes. This would require $O(n^2)$ operations, which is clearly too expensive. However, we can notice two things: First, the bounds of two nodes at the same level differ only by their degree. Therefore, for each level i , we can compute $2 \cdot \sum_{|j-i| \leq 1} \gamma_j + \sum_{|j-i| > 1} \gamma_j \cdot |j - i| - 2$ only once and then subtract $\deg(v)$ for each node at level i . We call the quantity $2 \cdot \sum_{|j-i| \leq 1} \gamma_j + \sum_{|j-i| > 1} \gamma_j \cdot |j - i| - 2$ the level-bound $L(i)$ of level i . Second, we can prove that $L(i)$ can actually be written as a function of $L(i-1)$.

LEMMA 7.2. *Let $L(i) := 2 \cdot \sum_{|j-i| \leq 1} \gamma_j + \sum_{|j-i| > 1} \gamma_j \cdot |j - i| - 2$. Also, let $\gamma_j = 0$ for $j \leq 0$ and $j > \max D$, where $\max D = \max_{v \in R(s)} d(s, v)$. Then $L(i) - L(i-1) = \sum_{j < i-2} \gamma_j - \sum_{j > i+1} \gamma_j$, $\forall i \in \{1, \dots, \max D\}$.*

PROOF. Since $\gamma_j = 0$ for $j \leq 0$ and $j > \max D$, we can write $L(i)$ as $2 \cdot (\gamma_{i-1} + \gamma_i + \gamma_{i+1}) + \sum_{|j-i| > 1} \gamma_j \cdot |j - i| - 2$, $\forall i \in \{1, \dots, \max D\}$. The difference between $L(i)$ and $L(i-1)$ is: $2 \cdot (\gamma_{i-1} + \gamma_i + \gamma_{i+1}) + \sum_{|j-i| > 1} |j - i| \cdot \gamma_j - 2 \cdot (\gamma_{i-2} + \gamma_{i-1} + \gamma_i) + \sum_{|j-i+1| > 1} |j - i + 1| \cdot \gamma_j = 2 \cdot (\gamma_{i+1} - \gamma_{i-2}) + 2 \cdot \gamma_{i-2} - 2 \cdot \gamma_{i+1} + \sum_{j < i-2 \cup j > i+1} (|j - i| - |j - i + 1|) \cdot \gamma_j = \sum_{j < i-2} \gamma_j - \sum_{j > i+1} \gamma_j$. \square

ALGORITHM 5: The updateBoundsLB function for undirected graphs

Input: A graph $G = (V, E)$, a source node s
Output: Lower bounds $L_s^{\text{LB}}(v, r(v))$ of each node $v \in R(s)$

```

1  $d \leftarrow \text{BFSfrom}(s)$ ;
2  $\max D \leftarrow \max_{v \in V} d(s, v)$ ;
3  $\text{sum}\Gamma_{\leq 0} \leftarrow 0$ ;  $\text{sum}\Gamma_{\leq -1} \leftarrow 0$ ;  $\text{sum}\Gamma_{> \max D + 1} \leftarrow 0$ ;
4 for  $i = 1, 2, \dots, \max D$  do
5    $\Gamma_i \leftarrow \{w \in V : d(s, w) = i\}$ ;
6    $\gamma_i \leftarrow \#\Gamma_i$ ;
7    $\text{sum}\Gamma_{\leq i} \leftarrow \text{sum}\Gamma_{\leq i-1} + \gamma_i$ ;
8    $\text{sum}\Gamma_{> i} \leftarrow |V| - \text{sum}\Gamma_{\leq i}$ ;
9  $L(1) \leftarrow \gamma_1 + \gamma_2 + \text{sum}\Gamma_{> 2} - 2$ ;
10 for  $i = 2, \dots, \max D$  do
11    $L(i) \leftarrow L(i-1) + \text{sum}\Gamma_{\leq i-3} - \text{sum}\Gamma_{> i+1}$ ;
12 for  $i = 1, \dots, \max D$  do
13   foreach  $v \in \Gamma_i$  do
14      $L_s^{\text{LB}}(v, r(v)) \leftarrow (L(i) - \deg(v)) \cdot \frac{(n-1)}{(r(v)-1)^2}$ ;
15 return  $L_s^{\text{LB}}(v, r(v)) \quad \forall v \in V$ 

```

Algorithm 5 describes the computation of $L_s^{\text{LB}}(v, r(v))$. First, we compute all the distances between s and the nodes in $R(s)$ with a BFS, storing the number of nodes in each level and the number of nodes in levels $j \leq i$ and $j > i$ respectively (Lines 1–8). Then, we compute the level bound $L(1)$ of level 1 according to its definition (Line 9) and those of the other level according to Lemma 7.2 (Line 11). The lower bound $L_s^{\text{LB}}(v, r(v))$ is then computed for each node v by subtracting its degree to $L(d(s, v))$ and normalizing (Line 14). The complexity of Lines 1–8 is that of running a BFS, i.e.,

$O(n + m)$. Line 11 is repeated once for each level (which cannot be more than n) and Line 14 is repeated once for each node in $R(s)$. Therefore, the following proposition holds.

PROPOSITION 7.3. *Computing the lower bound $L_s^{LB}(v, r(v))$ takes $O(n + m)$ time.*

For directed strongly connected graphs, the result does not hold for nodes w whose level is no greater than $d(s, v)$, since there might be a directed edge or a shortcut from v to w . Yet, for nodes w such that $d(s, w) > d(s, v)$, it is still true that $d(v, w) \geq d(s, w) - d(s, v)$. For the remaining nodes (apart from the outgoing neighbors of v), we can only say that the distance must be at least 2. The upper bound $L_s^{LB}(v, r(v))$ for directed graphs can therefore be defined as follows:

$$2 \cdot \#\{w \in R(s) : d(s, w) - d(s, v) \leq 1\} + \sum_{\substack{w \in R(s) \\ d(s, w) - d(s, v) > 1}} (d(s, w) - d(s, v)) - \deg(v) - 2. \quad (5)$$

The computation of $L_s^{LB}(v, r(v))$ for directed strongly-connected graphs is analogous to the one described in Algorithm 5.

8 THE DIRECTED DISCONNECTED CASE

In the directed disconnected case, even if the time complexity of computing strongly connected components is linear in the input size, the time complexity of computing the number of reachable vertices is much bigger (assuming SETH, it cannot be $O(m^{2-\epsilon})$ (Borassi 2016)). For this reason, when computing our upper bounds, we cannot rely on the exact value of $r(v)$: for now, let us assume that we know a lower bound $\alpha(v) \leq r(v)$ and an upper bound $\omega(v) \geq r(v)$. The definition of these bounds is postponed to Section 8.4.

Furthermore, let us assume that we have a lower bound $L(v, r(v))$ on the farness of v , depending on the number $r(v)$ of vertices reachable from v : in order to obtain a bound not depending on $r(v)$, the simplest approach is $f(v) \geq L(v, r(v)) \geq \min_{\alpha(v) \leq r \leq \omega(v)} L(v, r)$. However, during the algorithm, computing the minimum among all these values might be quite expensive, if $\omega(v) - \alpha(v)$ is big. In order to solve this issue, we find a small set $X \subseteq [\alpha(v), \omega(v)]$ such that $\min_{\alpha(v) \leq r \leq \omega(v)} L(v, r) = \min_{r \in X} L(v, r)$.

More specifically, we find a condition that is verified by “many” values of r , and that implies $L(v, r) \geq \min(L(v, r-1), L(v, r+1))$: this way, we may define X as the set of values of r that either do not verify this condition, or that are extremal points of the interval $[\alpha(v), \omega(v)]$ (indeed, all other values cannot be minima of $L(v, r)$). Since all our bounds are of the form $L(v, r) = \frac{(n-1)S(v, r)}{(r-1)^2}$, where $S(v, r)$ is a lower bound on $\sum_{w \in R(v)} d(v, w)$, we state our condition in terms of the function $S(v, r)$. For instance, in the case of the `updateBoundsBFSCut` function, $S_d^{\text{CUT}}(v, r) = \sum_{w \in N_d(v)} d(v, w) - \tilde{y}_{d+1}(v) + (d+2)(r - n_d(v))$, as in Lemma 6.1.

LEMMA 8.1. *Let v be a vertex, and let $S(v, r)$ be a positive function such that $S(v, r(v)) \leq \sum_{w \in R(v)} d(v, w)$ (where $r(v)$ is the number of vertices reachable from v). Assume that $S(v, r+1) - S(v, r) \leq S(v, r) - S(v, r-1)$. Then, if $L(v, r) := \frac{(n-1)S(v, r)}{(r-1)^2}$ is the corresponding bound on the farness of v , $\min(L(v, r+1), L(v, r-1)) \leq L(v, r)$.*

PROOF. Let us define $d = S(v, r+1) - S(v, r)$. Then, $L(v, r+1) \leq L(v, r)$ if and only if $\frac{(n-1)S(v, r+1)}{r^2} \leq \frac{(n-1)S(v, r)}{(r-1)^2}$ if and only if $\frac{S(v, r)+d}{r^2} \leq \frac{S(v, r)}{(r-1)^2}$ if and only if $(r-1)^2(S(v, r)+d) \leq r^2 S(v, r)$ if and only if $S(v, r)(r^2 - (r-1)^2) \geq (r-1)^2 d$ if and only if $S(v, r)(2r-1) \geq (r-1)^2 d$.

Similarly, if $d' = S(v, r) - S(v, r-1)$, $L(v, r-1) \leq L(v, r)$ if and only if $\frac{(n-1)S(v, r-1)}{(r-2)^2} \leq \frac{(n-1)S(v, r)}{(r-1)^2}$ if and only if $\frac{S(v, r)-d'}{(r-2)^2} \leq \frac{S(v, r)}{(r-1)^2}$ if and only if $(r-1)^2(S(v, r)-d') \leq (r-2)^2 S(v, r)$ if

and only if $S(v, r)((r-1)^2 - (r-2)^2) \leq (r-1)^2 d'$ if and only if $S(v, r)(2r-3) \leq (r-1)^2 d'$ if and only if $S(v, r)(2r-1) \leq (r-1)^2 d' + 2S(v, r)$.

We conclude that, assuming $d \leq d'$, $(r-1)^2 d \leq (r-1)^2 d' \leq (r-1)^2 d' + 2S(v, r)$, and one of the two previous conditions is always satisfied. \square

8.1 The Neighborhood-Based Lower Bound

In the neighborhood-based lower bound, we computed upper bounds $\tilde{y}_k(v)$ on $|\Gamma_k(v)|$, and we defined the lower bound $S^{\text{NB}}(v, r(v)) \leq \sum_{w \in R(v)} d(v, w)$, by

$$S^{\text{NB}}(v, r(v)) := \sum_{k=1}^{\text{diam}(G)} k \cdot \min \left\{ \tilde{y}_k(v), \max \left\{ r(v) - \sum_{i=0}^{k-1} \tilde{y}_i(v), 0 \right\} \right\}.$$

The corresponding bound on $f(v)$ is $L^{\text{NB}}(v, r(v)) := \frac{(n-1)S^{\text{NB}}(v, r(v))}{(r(v)-1)^2}$: let us apply Lemma 8.1 with $S(v, r) = S^{\text{NB}}(v, r)$ and $L(v, r) = L^{\text{NB}}(v, r)$. We have that the local minima of $L^{\text{NB}}(v, r(v))$ are obtained on values r such that $S^{\text{NB}}(v, r+1) - S^{\text{NB}}(v, r) > S^{\text{NB}}(v, r) - S^{\text{NB}}(v, r-1)$, that is, when $r = \sum_{i=0}^l \tilde{y}_i(v)$ for some l . Hence, our final bound $L^{\text{NB}}(v)$ becomes:

$$\min \left(L^{\text{NB}}(v, \alpha(v)), L^{\text{NB}}(v, \omega(v)), \min \left\{ L^{\text{NB}}(v, r) : \alpha(v) < r < \omega(v), r = \sum_{i=0}^l \tilde{y}_i(v) \right\} \right). \quad (6)$$

This bound can be computed with no overhead, by modifying Lines 20–25 in Algorithm 3. Indeed, when $r(v)$ is known, we have two cases: either $\text{nVisited}[s] < r(v)$, and we continue, or $\text{nVisited}[s] \geq r(v)$, and $S^{\text{NB}}(v, r(v))$ is computed. In the disconnected case, we need to distinguish three cases:

- (1) if $\text{nVisited}[v] < \alpha(v)$, we simply continue the computation;
- (2) if $\alpha(v) \leq \text{nVisited}[v] < \omega(v)$, we compute $L^{\text{NB}}(v, \text{nVisited}[v])$, and we update the minimum in Equation (6) (if this is the first occurrence of this situation, we also have to compute $L^{\text{NB}}(v, \alpha(v))$);
- (3) if $\text{nVisited}[v] \geq \omega(v)$, we compute $L^{\text{NB}}(v, \omega(v))$, and we update the minimum in Equation (6).

Since this procedure needs time $O(1)$, it has no impact on the running time of the computation of the neighborhood-based lower bound.

8.2 The updateBoundsBFSCut Function

Let us apply Lemma 8.1 to the bound used in the updateBoundsBFSCut function. In this case, by Lemma 6.1, $S_d^{\text{CUT}}(v, r) = \sum_{w \in N_d(v)} d(v, w) - \tilde{y}_{d+1}(v) + (d+2)(r - n_d(v))$, and $S_d^{\text{CUT}}(v, r+1) - S_d^{\text{CUT}}(v, r) = d+2$, which does not depend on r . Hence, the condition in Lemma 8.1 is always verified, and the only values we have to analyze are $\alpha(v)$ and $\omega(v)$. Hence, the lower bound becomes $f(v) \geq L_d^{\text{CUT}}(v, r(v)) \geq \min_{\alpha(v) \leq r \leq \omega(v)} L_d^{\text{CUT}}(v, r) = \min(L_d^{\text{CUT}}(v, \alpha(v)), L_d^{\text{CUT}}(v, \omega(v)))$ (which does not depend on $r(v)$).

This means that, in order to adapt the updateBoundsBFSCut function (Algorithm 4), it is enough to replace Lines 10, 22 in order to compute both $L_d^{\text{CUT}}(v, \alpha(v))$ and $L_d^{\text{CUT}}(v, \omega(v))$, and to replace Lines 11, 23 in order to stop if $\min(L_d^{\text{CUT}}(v, \alpha(v)), L_d^{\text{CUT}}(v, \omega(v))) \geq x$.

8.3 The updateBoundsLB Function

In this case, we do not apply Lemma 8.1 to obtain simpler bounds. Indeed, the updateBoundsLB function improves the bounds of vertices that are quite close to the source of the BFS, and hence

are likely to be in the same component as this vertex. Consequently, if we perform a BFS from a vertex s , we can simply compute $L_s^{\text{LB}}(v, r(v))$ for all vertices in the same strongly connected component as s , and for these vertices we know the value $r(v) = r(s)$. The computation of better bounds for other vertices is left as an open problem.

8.4 Computing $\alpha(v)$ and $\omega(v)$

It now remains to compute $\alpha(v)$ and $\omega(v)$. This can be done during the preprocessing phase of our algorithm, in linear time. To this purpose, let us precisely define the node-weighted directed acyclic graph $\mathcal{G} = (\mathcal{V}, \mathcal{E})$ of strongly connected components (in short, SCCs) corresponding to a directed graph $G = (V, E)$. In this graph, \mathcal{V} is the set of SCCs of G , and, for any two SCCs $C, D \in \mathcal{V}$, $(C, D) \in \mathcal{E}$ if and only if there is an arc in E from a node in C to a node in D . For each SCC $C \in \mathcal{V}$, the weight $w(C)$ of C is equal to $|C|$, that is, the number of nodes in the SCC C . Note that the graph \mathcal{G} is computable in linear time.

For each node $v \in C$, $r(v) = \sum_{D \in R(C)} w(D)$, where $R(C)$ denotes the set of SCCs that are reachable from C in \mathcal{G} . This means that we simply need to compute a lower (respectively, upper) bound $\alpha_{\text{SCC}}(C)$ (respectively, $\omega_{\text{SCC}}(C)$) on $\sum_{D \in R(C)} w(D)$, for every SCC C . To this aim, we first compute a topological sort $\{C_1, \dots, C_l\}$ of \mathcal{V} (that is, if $(C_i, C_j) \in \mathcal{E}$, then $i < j$). Successively, we use a dynamic programming approach, and, by starting from C_l , we process the SCCs in reverse topological order, and we set the following:

$$\alpha_{\text{SCC}}(C) = w(C) + \max_{(C,D) \in \mathcal{E}} \alpha_{\text{SCC}}(D) \quad \omega_{\text{SCC}}(C) = w(C) + \sum_{(C,D) \in \mathcal{E}} \omega_{\text{SCC}}(D).$$

Note that processing the SCCs in reverse topological ordering ensures that the values $\alpha(D)$ and $\omega(D)$ on the right hand side of these equalities are available when we process the SCC C . Clearly, the complexity of computing $\alpha(C)$ and $\omega(C)$, for each SCC C , is linear in the size of \mathcal{G} , which in turn is smaller than G .

Observe that the bounds obtained through this simple approach can be improved by using some “tricks.” First of all, when the biggest SCC \tilde{C} is processed, we do not use the dynamic programming approach and we exactly compute $\sum_{D \in R(\tilde{C})} w(D)$ by performing a BFS starting from any node in \tilde{C} . This way, not only $\alpha(\tilde{C})$ and $\omega(\tilde{C})$ are exact, but also $\alpha_{\text{SCC}}(C)$ and $\omega_{\text{SCC}}(C)$ are improved for each SCC C from which it is possible to reach \tilde{C} . Finally, in order to compute the upper bounds for the SCCs that are able to reach \tilde{C} , we can run the dynamic programming algorithm on the graph obtained from \mathcal{G} by removing all components reachable from \tilde{C} , and we can then add $\sum_{D \in R(\tilde{C})} w(D)$.

The pseudocode is available in Algorithm 6. The complexity of this algorithm is linear in the number of edges of G . Indeed, computing the graph of strongly connected components can be done in time linear in E , computing the components reachable from \tilde{C} and the components that can reach \tilde{C} can be done in time linear in \mathcal{E} , and executing the dynamic programming procedure requires time linear in \mathcal{E} (since each edge of \mathcal{G} is analysed only once).

9 EXPERIMENTAL RESULTS

In this section, we test the four variations of our algorithm on several real-world networks, in order to evaluate their performances. All the networks used in our experiments come from the datasets SNAP (snap.stanford.edu/), NEXUS (igraph.org/r/doc/nexus.html), LASAGNE (lasagne-unifi.sourceforge.net/), LAW (law.di.unimi.it), KONECT (konect.uni-koblenz.de), and IMDB (www.imdb.com). The platform for our tests is a shared-memory server with 256GB RAM and 2x8 Intel(R) Xeon(R) E5-2680 cores (32 threads due to hyperthreading) at 2.7GHz. The algorithms are implemented in C++, building on the open-source *NetworKit* framework (Staudt et al. 2016).

ALGORITHM 6: Estimating the number of reachable vertices in directed disconnected graphs

Input: A graph $G = (V, E)$
Output: Lower and upper bounds $\alpha(v)$, $\omega(v)$ on the number of vertices reachable from v

```

1   $(\mathcal{V}, \mathcal{E}, w) \leftarrow \text{computeSCCGraph}(G)$ ;
2   $\tilde{C} \leftarrow$  the biggest SCC;
3   $\alpha_{\text{SCC}}(\tilde{C}), \omega_{\text{SCC}}(\tilde{C}) \leftarrow$  the number of vertices reachable from  $\tilde{C}$ ;
4  for  $X \in \mathcal{V}$  in reverse topological order do
5      if  $X == \tilde{C}$  then continue;
6       $\alpha_{\text{SCC}}(X), \omega_{\text{SCC}}(X), \omega'_{\text{SCC}}(X) \leftarrow 0$ ;
7      for  $Y$  neighbor of  $X$  in  $\mathcal{G}$  do
8           $\alpha_{\text{SCC}}(X) \leftarrow \max(\alpha_{\text{SCC}}(X), \alpha_{\text{SCC}}(Y))$ ;
9           $\omega_{\text{SCC}}(X) \leftarrow \omega_{\text{SCC}}(X) + \omega_{\text{SCC}}(Y)$ ;
10         if  $Y$  not reachable from  $\tilde{C}$  then  $\omega'_{\text{SCC}}(X) \leftarrow \omega'_{\text{SCC}}(X) + \omega_{\text{SCC}}(Y)$ ;
11     if  $X$  reaches  $\tilde{C}$  then  $\omega_{\text{SCC}}(X) \leftarrow \omega'_{\text{SCC}}(X) + \omega_{\text{SCC}}(\tilde{C})$ ;
12      $\alpha_{\text{SCC}}(X) \leftarrow \alpha_{\text{SCC}}(X) + w(X)$ ;
13      $\omega_{\text{SCC}}(X) \leftarrow \omega_{\text{SCC}}(X) + w(X)$ ;
14 for  $v \in V$  do
15      $\alpha(v) = \alpha_{\text{SCC}}(\text{the component of } v)$ ;
16      $\omega(v) = \omega_{\text{SCC}}(\text{the component of } v)$ ;
17 return  $\alpha, \omega$ 

```

9.1 Comparison with the State of the Art

In order to compare the performance of our algorithm with state-of-the-art approaches, we select 19 directed complex networks, 17 undirected complex networks, 6 directed street networks, and 6 undirected street networks (the undirected versions of the previous ones). The number of nodes of most of these networks ranges between 5,000 and 100,000. We test four different variations of our algorithm, that provide different implementations of the functions `computeBounds` and `updateBounds` (for more information, we refer to Section 4):

- (1) `DEGCUT` uses the conservative strategies `computeBoundsDeg` and `updateBoundsBFSCut`;
- (2) `DEGBOUND` uses the conservative strategy `computeBoundsDeg` and the aggressive strategy `updateBoundsLB`;
- (3) `NBCUT` uses the aggressive strategy `computeBoundsNB` and the conservative strategy `updateBoundsBFSCut`;
- (4) `NBBOUND` uses the aggressive strategies `computeBoundsNB` and `updateBoundsLB`.

We compare these algorithms with our implementations of the best existing algorithms for top- k closeness centrality.⁴ The first one (Olsen et al. 2014) is based on a pruning technique and on Δ -BFS, a method to reuse information collected during a BFS from a node to speed up a BFS from one of its in-neighbors; we denote this algorithm as `OLH`. The second one, `OCL` (Okamoto et al. 2008), performs some BFSes from a random sample of nodes to estimate the closeness centrality of all the other nodes, then it computes the exact centrality of all the nodes whose estimate is big enough. Note that this algorithm requires the input graph to be (strongly) connected: for this reason, differently from the other algorithms, we have run this algorithm on the largest (strongly) connected component of the input graph. Also, we recall that the algorithms we propose, as well as `OLH` are *exact algorithms*, i.e., they all return the top- k nodes with highest closeness centrality

⁴Note that the source code of our competitors is not available.

Table 2. Complex Networks: Geometric Mean and Standard Deviation of the Edge Traversal Ratios of the Algorithm in (Olsen et al. 2014) (OLH), the Algorithm in (Okamoto et al. 2008) (OCL), and the Four Variations of the New Algorithm (DEGCUT, DEGBOUND, NBCUT, NBBOUND)

k	ALGORITHM	DIRECTED		UNDIRECTED		BOTH	
		GMEAN	GSTDDEV	GMEAN	GSTDDEV	GMEAN	GSTDDEV
1	OLH	21.24	5.68	11.11	2.91	15.64	4.46
	OCL	1.71	1.54	2.71	1.50	2.12	1.61
	DEGCUT	104.20	6.36	171.77	6.17	131.94	6.38
	DEGBOUND	3.61	3.50	5.83	8.09	4.53	5.57
	NBCUT	123.46	7.94	257.81	8.54	174.79	8.49
	NBBOUND	17.95	10.73	56.16	9.39	30.76	10.81
10	OLH	21.06	5.65	11.11	2.90	15.57	4.44
	OCL	1.31	1.31	1.47	1.11	1.38	1.24
	DEGCUT	56.47	5.10	60.25	4.88	58.22	5.00
	DEGBOUND	2.87	3.45	2.04	1.45	2.44	2.59
	NBCUT	58.81	5.65	62.93	5.01	60.72	5.34
	NBBOUND	9.28	6.29	10.95	3.76	10.03	5.05
100	OLH	20.94	5.63	11.11	2.90	15.52	4.43
	OCL	1.30	1.31	1.46	1.11	1.37	1.24
	DEGCUT	22.88	4.70	15.13	3.74	18.82	4.30
	DEGBOUND	2.56	3.44	1.67	1.36	2.09	2.57
	NBCUT	23.93	4.83	15.98	3.89	19.78	4.44
	NBBOUND	4.87	4.01	4.18	2.46	4.53	3.28

Bold values represent better higher values.

along with their exact scores. OCL is exact with high probability (and returned the exact ranking in all our experiments).

In order to perform a fair comparison, we consider the *edge traversal ratio*, which is defined as $\frac{mn}{m_{vis}}$ in directed graphs, $\frac{2mn}{m_{vis}}$ in undirected graphs, where m_{vis} is the number of arcs visited during the algorithm, and mn (respectively, $2mn$) is an estimate of the number of arcs visited by the *textbook* algorithm in directed (resp., undirected) graphs. For OCL, the edge traversal ratio is computed using the number of nodes and edges in the largest (strongly) connected component of the graph. Note that the edge traversal ratio does not depend on the implementation, nor on the machine used for the algorithm, and it does not consider parts of the code that need subquadratic time in the worst case. These parts are negligible in our algorithm, because their worst case running time is $O(n \log n)$ or $O(mD)$ where D is the diameter of the graph, but they can be significant when considering the competitors. For instance, in the particular case of OLH, we have just counted the arcs visited in BFS and Δ -BFS, ignoring all the operations done in the pruning phases (see (Olsen et al. 2014)).

We consider the geometric mean of the edge traversal ratios over all graphs in the dataset. In our opinion, this quantity is more informative than the arithmetic mean, which is highly influenced by the maximum value. Our choice is further confirmed by the geometric standard deviation, which is always quite small.

The results are summarized in Table 2 for complex networks and Table 3 for street networks. For references of the edge traversal ratios of each single graph, we refer to Tables 9–11 in Appendix B. Our evaluation focuses on an interpretation of aggregated results, though.

On complex networks, the best algorithm is NBCUT: when $k = 1$, the edge traversal ratios are always bigger than 100, up to 258, when $k = 10$ they are close to 60, and when $k = 100$ they are

Table 3. Street Networks: Geometric Mean and Standard Deviation of the Edge Traversal Ratios of the Algorithm in (Olsen et al. 2014) (OLH), the Algorithm in (Okamoto et al. 2008) (OCL), and the Four Variations of the New Algorithm (DEGCUT, DEGBOUND, NBCUT, NBBOUND)

k	ALGORITHM	DIRECTED		UNDIRECTED		BOTH	
		GMEAN	GSTDDEV	GMEAN	GSTDDEV	GMEAN	GSTDDEV
1	OLH	4.11	1.83	4.36	2.18	4.23	2.01
	OCL	3.39	1.28	3.23	1.28	3.31	1.28
	DEGCUT	4.14	2.07	4.06	2.06	4.10	2.07
	DEGBOUND	187.10	1.65	272.22	1.67	225.69	1.72
	NBCUT	4.12	2.07	4.00	2.07	4.06	2.07
	NBBOUND	250.66	1.71	382.47	1.63	309.63	1.74
10	OLH	4.04	1.83	4.28	2.18	4.16	2.01
	OCL	2.93	1.24	2.81	1.24	2.87	1.24
	DEGCUT	4.09	2.07	4.01	2.06	4.05	2.07
	DEGBOUND	172.06	1.65	245.96	1.68	205.72	1.72
	NBCUT	4.08	2.07	3.96	2.07	4.02	2.07
	NBBOUND	225.26	1.71	336.47	1.68	275.31	1.76
100	OLH	4.03	1.82	4.27	2.18	4.15	2.01
	OCL	2.90	1.24	2.79	1.24	2.85	1.24
	DEGCUT	3.91	2.07	3.84	2.07	3.87	2.07
	DEGBOUND	123.91	1.56	164.65	1.67	142.84	1.65
	NBCUT	3.92	2.08	3.80	2.09	3.86	2.08
	NBBOUND	149.02	1.59	201.42	1.69	173.25	1.67

Bold values represent better higher values.

close to 20. Another good option is DEGCUT, which achieves results similar to NBCUT, but it has almost no overhead at the beginning (while NBCUT needs a preprocessing phase with cost $O(mD)$). Furthermore, DEGCUT is very easy to implement, becoming a very good candidate for state-of-the-art graph libraries. The edge traversal ratios of the competitors are smaller: OLH has edge traversal ratios between 10 and 20, and OCL provides almost no improvement with respect to the *textbook* algorithm.

On street networks, the best option is NBBOUND: for $k = 1$, the average improvement is about 250 in the directed case and about 382 in the undirected case, and it always remains bigger than 150, even for $k = 100$. It is worth noting that also the performance of DEGBOUND are quite good, being at least 70% of NBBOUND. Even in this case, the DEGBOUND algorithm offers some advantages: it is very easy to be implemented, and there is no overhead in the first part of the computation. All the competitors perform relatively poorly on street networks, since their improvement is always smaller than 5.

Tables 6, 7, and 8 compare the running times of the different algorithms, for $k = 1$, $k = 10$, and $k = 100$, respectively. All times refer to a sequential implementation of the different algorithms. Also, we recall that the running times of our algorithms and OLH refer to the whole graphs, whereas OCL is run on the largest (strongly) connected component (SCC). Since for some graphs this only contains a small fraction of the total number of nodes and edges, sometimes the running time of OCL is smaller than those of our algorithms. In particular, this is the case for five directed graphs: `out.subelj_jung-j_jung-j`, `out.subelj_jdk_jdk`, `out.ego-twitter`, `out.ego-gplus`, and `out.subelj_coracora`. However, notice that in the first four graphs, the number of nodes in the largest SCC is less than 1% of the total number of nodes. For `out.subelj_coracora`, the nodes in the largest SCC are about 17% of the total number of nodes. When run on the largest SCC

of the graphs, our algorithms always took less time than OCL. For all other graphs, at least one of our four algorithms is faster than both OLH and OCL, in most cases by a wide margin.

Overall, we conclude that the preprocessing function `computeBoundsNB` always leads to better results (in terms of visited edges) than `computeBoundsDeg`, but the difference is quite small. Also in terms of running times the differences are mostly small, and `computeBoundsDeg` sometimes is even faster. Hence, in some cases, `computeBoundsDeg` could be even preferred, because of its simplicity. Conversely, the performance of `updateBoundsBFSCut` is very different from the performance of `updateBoundsLB`: the former works much better on complex networks, while the latter works much better on street networks. Currently, these two approaches exclude each other: an open problem left by this work is the design of a “combination” of the two, which works both in complex networks and in street networks. Finally, the experiments show that the best variation of our algorithm outperforms all competitors in all frameworks considered: both in complex and in street networks, both in directed and undirected graphs.

Harmonic centrality. As mentioned in the introduction, all our methods can be easily generalized to any centrality measure in the form $c(v) = \sum_{w \neq v} f(d(v, w))$, where f is a decreasing function such that $f(+\infty) = 0$. We also implemented a version of `DEGCUT`, `DEGBOUND`, `NBCUT`, and `NBBOUND` for *harmonic centrality*, which is defined as $h(v) = \sum_{w \neq v} \frac{1}{d(v, w)}$. Also for harmonic centrality, we compute the edge traversal ratios on the textbook algorithm.

For the complex networks used in our experiments, finding the k nodes with highest harmonic centrality is *always faster* than finding the k nodes with highest closeness, for all four methods and k values in $\{1, 10, 100\}$. For example, for `NBCUT` and $k = 1$, the geometric mean⁵ of the edge traversal ratios is 486.07, whereas for closeness it is 174.79 (as reported in Table 2).

For street networks, the version of harmonic centrality is faster than the version for closeness for `DEGCUT` and `NBCUT`, but it is slower for `DEGBOUND` and `NBBOUND`. In particular, the average (geometric mean) edge traversal ratio of `NBBOUND` for harmonic centrality is 103.58 for $k = 1$, 93.49 for $k = 10$, and 62.22 for $k = 100$, which is about a factor 3 smaller than the edge traversal ratio of `NBBOUND` for closeness (see Table 3). Nevertheless, this is significantly faster than the textbook algorithm.

9.2 Real-World Large Networks

In this section, we run our algorithm on bigger inputs, by considering a dataset containing 23 directed networks, 15 undirected networks, and 5 street networks, with up to 3,774,768 nodes and 117,185,083 edges. On this dataset, we run the fastest variant of our algorithm (`DEGBOUND` in complex networks, `NBBOUND` in street networks), using 64 threads (however, the server used has only 16 cores and runs 32 threads with hyperthreading; we account for memory latency in graph computations by oversubscribing slightly).

Once again, we consider the *edge traversal ratio*, which is defined as $\frac{mn}{m_{vis}}$ in directed graphs, $\frac{2mn}{m_{vis}}$ in undirected graphs. It is worth observing that we are able to compute for the first time the k most central nodes (with respect to the closeness measure) of networks with millions of nodes and hundreds of millions of arcs, with $k = 1$, $k = 10$, and $k = 100$. The detailed results are shown in Table 12 in Appendix C, where for each network we report the running time (in format h:mm:ss) and the edge traversal ratio. A summary of these results is available in Table 4, which contains the geometric means of the edge traversal ratios, with the corresponding standard deviations.

For $k = 1$, the geometric mean of the edge traversal ratios is always above 200 in complex networks, and above 700 in street networks. In undirected graphs, the edge traversal ratios are even

⁵We report the geometric mean over both directed and undirected networks.

Table 4. Big Networks: Geometric Mean and Standard Deviation of the Edge Traversal Ratios of the Best Variation of the New Algorithm (DEGBOUND in Complex Networks, NBBBOUND in Street Networks)

Input	k	DIRECTED		UNDIRECTED		BOTH	
		GMEAN	GSTDDEV	GMEAN	GSTDDEV	GMEAN	GSTDDEV
Street	1	742.42	2.60	1681.93	2.88	1117.46	2.97
	10	724.72	2.67	1673.41	2.92	1101.25	3.03
	100	686.32	2.76	1566.72	3.04	1036.95	3.13
Complex	1	247.65	11.92	551.51	10.68	339.70	11.78
	10	117.45	9.72	115.30	4.87	116.59	7.62
	100	59.96	8.13	49.01	2.93	55.37	5.86

bigger: close to 500 in complex networks and close to 1,600 in street networks. For bigger values of k , the performance does not decrease significantly: on complex networks, the edge traversal ratios are bigger than or very close to 50, even for $k = 100$. In street networks, the performance loss is even smaller, always below 10% for $k = 100$.

Regarding the robustness of the algorithm, we outline that the algorithm always achieves performance improvements bigger than \sqrt{n} in street networks, and that in complex networks, with $k = 1$, 64% of the networks have edge traversal ratios above 100, and 33% of the networks above 1,000. In some cases, the edge traversal ratio is even bigger: in the com-Orkut network, our algorithm for $k = 1$ is almost 35,000 times faster than the *textbook* algorithm.

In our experiments, we also report the running time of our algorithm. Even for $k = 100$, a few minutes are sufficient to conclude the computation on most networks, and, in all but two cases, the total time is smaller than 3 hours. For $k = 1$, the computation always terminates in at most 1 hour and a half, apart from two street networks where it needs less than 2 hours and a half. Overall, the total time needed to compute the most central vertex in all the networks is smaller than 1 day. In contrast to this, if we extrapolate the results in Tables 2 and 3, it seems plausible that the fastest competitor OLH would require a month or so.

10 IMDB CASE STUDY

In this section, we apply the new algorithm NBCUT to analyse the IMDB graph, where nodes are actors, and two actors are connected if they played together in a movie (TV-series are ignored). The data collected comes from the website <http://www.imdb.com>. In line with <http://oracleofbacon.org>, we decide to exclude some genres from our database: awards-shows, documentaries, game-shows, news, realities and talk-shows. We analyze snapshots of the actor graph, taken every 5 years from 1940 to 2010, and 2014. The results are reported in Tables 13 and 14 in the Appendix.

The algorithm. Thanks to this experiment, we can evaluate the performance of our algorithm on increasing snapshots of the same graph. This way, we can have an informal idea on the asymptotic behavior of its complexity. In Figure 3, we have plotted the edge traversal ratio with respect to the number of nodes: if the edge traversal ratio is I , the running time is $O(\frac{mn}{I})$. Hence, assuming that $I = cn$ for some constant c (which is approximately verified in the actor graph, as shown by Figure 3), the running time is linear in the input size. The total time needed to perform the computation on all snapshots is little more than 30 minutes for $k = 1$, and little more than 45 minutes for $k = 10$.

The results. In 2014, the most central actor is Michael Madsen, whose career spans 25 years and more than 170 films. Among his most famous appearances, he played as *Jimmy Lennox* in *Thelma*




Fig. 3. Growth of performance ratio with respect to the number of nodes ($k = 1$).

& Louise (Ridley Scott, 1991), as *Glen Greenwood* in *Free Willy* (Simon Wincer, 1993), as *Bob* in *Sin City* (Frank Miller, Robert Rodriguez, Quentin Tarantino), and as *Deadly Viper Budd* in *Kill Bill* (Quentin Tarantino, 2003–2004). The second is Danny Trejo, whose most famous movies are *Heat* (Michael Mann, 1995), where he played as *Trejo*, *Machete* (Ethan Maniquis, Robert Rodriguez, 2010) and *Machete Kills* (Robert Rodriguez, 2013), where he played as *Machete*. The third “actor” is not really an actor: he is the German dictator Adolf Hitler (he was also the most central actor in 2005 and 2010, and he was in the top 10 since 1990). This is a consequence of his appearances in several archive footages, that were re-used in several movies (he counts 775 credits, even if most of them are in documentaries or TV shows, which were eliminated). Among the movies where Adolf Hitler is credited, we find *Zelig* (Woody Allen, 1983), and *The Imitation Game* (Morten Tyldum, 2014). Among the other most central actors, we find many people who played a lot of movies, and most of them are quite important actors. However, this ranking does not discriminate between important roles and marginal roles: for instance, the actress Bess Flowers is not widely known, because she rarely played significant roles, but she appeared in over 700 movies in her 41 years career, and for this reason she was the most central for 30 years, between 1950 and 1980. Finally, it is worth noting that we never find Kevin Bacon in the top 10, even if he became famous for the “Six Degrees of Kevin Bacon” game (<http://oracleofbacon.org>). In this game the player receives an actor x and has to find a path of length at most 6 from x to Kevin Bacon in the actor graph. Kevin Bacon was chosen as the goal because he played in several movies, and he was thought to be one of the most central actors: this work shows that, actually, he is quite far from the top. Indeed, his closeness centrality is 0.336, while the most central actor has centrality 0.354, the 10th actor has centrality 0.350, and the 100th actor has centrality 0.341.

11 WIKIPEDIA CASE STUDY

In this section, we apply the new algorithm NBCUT to analyze the Wikipedia graph, where nodes are pages, and there is a directed edge from page p to page q if p contains a link to q . The data collected comes from DBPedia 3.7 (wiki.dbpedia.org). We analyze both the standard graph and the reverse graph, which contains an edge from page p to page q if q contains a link to p . The 10 most central pages are available in Table 5.

The algorithm. In the standard graph, the edge traversal ratio is 1,784 for $k = 1$, 1,509 for $k = 10$, and 870 for $k = 100$. The total running time is about 39 minutes for $k = 1$, 45 minutes for $k = 10$, and less than 1 hour and 20 minutes for $k = 100$. In the reversed graph, the algorithm performs even better: the edge traversal ratio is 87,918 for $k = 1$, 71,923 for $k = 10$, and 21,989 for $k = 100$. The total running times are less than 3 minutes for both $k = 1$ and $k = 10$, and less than 10 minutes for $k = 100$.

Table 5. Top 10 Pages in Wikipedia Directed Graph, Both in the Standard Graph and in the Reversed Graph

Position	Standard Graph	Reversed Graph
1ST	1989	United States
2ND	1967	World War II
3RD	1979	United Kingdom
4TH	1990	France
5TH	1970	Germany
6TH	1991	English language
7TH	1971	Association football
8TH	1976	China
9TH	1945	World War I
10TH	1965	Latin

The results. If we consider the standard graph, the results are quite unexpected: indeed, all the most central pages are years (the first is 1989). However, this is less surprising if we consider that these pages contain a lot of links to events that happened in that year. For instance, the out-degree of 1989 is 1,560, and the links contain pages from very different topics: historical events, like the fall of Berlin wall, days of the year, different countries where particular events happened, and so on. A similar argument also works for other years: indeed, the second page is 1967 (with out-degree 1,438), and the third is 1979 (with out-degree 1,452). Furthermore, all the 10 most central pages have out-degree at least 1,269. Overall, we conclude that the central page in the Wikipedia standard graph are not the “intuitively important” pages, but they are the pages that have a biggest number of links to pages with different topics, and this maximum is achieved by pages related to years.

Conversely, if we consider the reversed graph, the most central page is *United States*, confirming a common conjecture. Indeed, in <http://wikirank.di.unimi.it/>, it is shown that the United States are the center according to harmonic centrality, and many other measures (however, in that work, the ranking is only approximated). A further evidence for this conjecture comes from the Six Degree of Wikipedia game (<http://thewikigame.com/6-degrees-of-wikipedia>), where a player is asked to go from one page to the other following the smallest possible number of link: a hard variant of this game forces the player not to pass the *United States* page, which is considered to be central. In this work, we show that this conjecture is true. The second page is *World War II*, and the third is *United Kingdom*, in line with the results obtained by other centrality measures (see <http://wikirank.di.unimi.it/>), especially for the first two pages.

Overall, we conclude that most of the central pages in the reversed graph are nations, and that the results capture our intuitive notion of “important” pages in Wikipedia. Thanks to this new algorithm, we can compute these pages in a bit more than 1 hour for the original graph, and less than 10 minutes for the reversed one.

12 CONCLUSIONS

In this article, we have presented a hardness result on the computation of the most central vertex in a graph, according to closeness centrality. Then, we have presented a very efficient algorithm for the exact computation of the k most central vertices, presenting four different variants. Even if the time complexity of the new algorithm is equal to the time complexity of the textbook algorithm (which, in any case, cannot be improved in general), we have shown that in practice the former improves the latter by several orders of magnitude. In this regard, notice that our algorithm is specifically designed for scenarios where only the very central nodes are needed. In fact, preliminary

results show that for larger values of k (e.g., $k = 10^3$ and $k = 10^4$), the running times become closer to those of the textbook algorithm. In fact, it is important to notice that the speedup can never be larger than n/k , since at least k BFSs are needed to compute the closeness of the top- k nodes.

Our results also show that the variant NBCUT performs best on complex networks, whereas the variant NBBOUND performs best on networks with large diameter, such as street networks. We conjecture that this stems from NBCUT taking neighborhood sizes (among others) into account and NBBOUND distance differences: neighborhood sizes differ more in complex networks and distances differ more in high-diameter networks. The choice of which variant of our algorithm should be used therefore depends on the properties of the network under consideration. Determining which properties should be taken into account and whether our conjecture above can be verified are very interesting (theoretical) open questions.

We have also shown that the new algorithm outperforms the state of the art (whose time complexity is still equal to the complexity of the textbook algorithm), and we have computed for the first time the most central nodes (with respect to the closeness measure) in networks with millions of nodes and hundreds of millions of edges. Finally, we have considered as a case study several snapshots of the IMDB actor network, and the Wikipedia graph.

APPENDICES

A PROOF OF HARDNESS RESULT

This section is devoted to the proof of Theorem 3.1. We construct a reduction from the l -TwoDisjointSet problem, that is, finding two disjoint sets in a collection C of subsets of a given ground set X , where $|X| = O(\log^l(|C|))$. For example, X could be the set of numbers between 0 and h , and C could be the collection of subsets of even numbers between 0 and h (in this case, the answer is True, since there are two disjoint sets in the collection). It is simple to prove that this problem is equivalent to the Orthogonal Vector problem, by replacing a set X with its characteristic vector in $\{0, 1\}^{|X|}$ (Borassi et al. 2015a); consequently, an algorithm solving this problem in $O(|C|^{2-\epsilon})$ would falsify the Orthogonal Vector conjecture. For a direct reduction between the l -TwoDisjointSet problem and SETH, we refer to (Williams 2005) (where the TwoDisjointSet problem is named COOPERATIVESUBSETQUERY).

Given an instance (X, C) of the l -TwoDisjointSet problem, and given a set $C \in \mathcal{C}$, let R_C be $|\{C' \in C : C \cap C' \neq \emptyset\}|$. The TwoDisjointSet problem has no solutions if and only if $R_C = |C|$ for all $C \in \mathcal{C}$; indeed, $R_C = |C|$ means that C nontrivially intersects all the sets in C . We construct a directed graph $G = (V, E)$, where $|V|, |E| = O(|C||X|) = O(|C| \log^l |C|)$, such that

- (1) V contains a set of vertices C_0 representing the sets in C (from now on, if $C \in \mathcal{C}$, we denote by C_0 the corresponding vertex in C_0);
- (2) the centrality of C_0 is a function $c(R_C)$, depending only on R_C (that is, if $R_C = R_{C'}$ then $c(C_0) = c(C'_0)$);
- (3) the function $c(R_C)$ is decreasing with respect to R_C ;
- (4) the most central vertex is in C_0 .

In such a graph, the vertex with maximum closeness corresponds to the set S minimizing R_S : indeed, it is in C_0 by Condition 4, and it minimizes R_S by Conditions 2–3. Hence, assuming we can find S_0 in time $O(n^{2-\epsilon})$, we can easily check if the closeness of S_0 is $c(|C|)$: if it is not, it means that the corresponding TwoDisjointSet instance has a solution of the form (S, S_1) because $R_S \neq |C|$. Otherwise, for each C , $R_C \geq R_S = |C|$, because $c(C_0) \leq c(S_0) = c(|C|)$, and c is decreasing with respect to R_C . This means that $R_C = |C|$ for each C , and there are no two disjoint sets. This way, we can solve the l -TwoDisjointSet problem in $O(n^{2-\epsilon}) = O((|C| \log^l |C|)^{2-\epsilon}) = O(|C|^{2-\frac{\epsilon}{l}})$,




Fig. 4. Reducing the TwoDisjointSet problem to the problem of finding the most closeness central vertex.

against the Orthogonal Vector conjecture, and SETH. If we also want the graph to be sparse, we can add $O(|C| \log^l |C|)$ nodes with no outgoing edge.

Remark A.1. One might also want the graph in the reduction to be weakly connected, a condition that would be broken by the addition of these new nodes. To fix this, we can also add one edge from each of the newly added nodes to a node in Z : this way, the closeness of all the newly added vertices would be $\frac{1}{n-1}$, which is smaller than the closeness of all the vertices in C_0 .

To construct this graph (see Figure 4), we start by adding to V the copy C_0 of C , another copy C_1 of C , and a copy X_1 of X . These vertices are connected as follows: for each element $x \in X$ and set $C \in C$, we add an edge (C_0, x) and (x, C_1) , where C_0 is the copy of C in C_0 , and C_1 is the copy of C in C_1 . Moreover, we add a copy X_2 of X and we connect all pairs (C_0, x) with $C \in C$, $x \in X$ and $x \notin C$. This way, the closeness centrality of a vertex $C_0 \in C_0$ is $\frac{(|X|+R_C)^2}{(n-1)(|X|+2R_C)}$ (which only depends on R_C). To enforce Conditions 3–4, we add a path of length p leaving each vertex in C_1 , and q vertices linked to each vertex in C_0 , each of which has out-degree $|C|$: we show that by setting $p = 7$ and $q = 36$, all required conditions are satisfied.

More formally, we have constructed the following graph $G = (V, E)$:

- $V = Z \cup Y \cup C_0 \cup X_1 \cup X_2 \cup C_1 \cup \dots \cup C_p$, where Z is a set of cardinality $q|C|$, Y a set of cardinality q , the C_i s are copies of C and the X_i s are copies of X ;
- each vertex in Y has $|C|$ neighbors in Z , and these neighbors are disjoint;
- for each $x \in C$, there are edges from $C_0 \in C_0$ to $x \in X_1$, and from $x \in X_1$ to $C_1 \in C_1$;
- for each $x \notin C$, there is an edge from $C_0 \in C_0$ to $x \in X_2$;
- each $C_i \in C_i$, $1 \leq i \leq p$, is connected to the same set $C_{i+1} \in C_{i+1}$;
- no other edge is present in the graph.

Note that the number of edges in this graph is $O(|C||X|) = O(|C| \log^l(|C|))$, because $|X| < \log^l(|C|)$,

LEMMA A.2. *Assuming $|C| > 1$, all vertices outside C_0 have closeness centrality at most $\frac{2|C|}{n-1}$, where n is the number of vertices.*

PROOF. If a vertex is in Z, X_2 , or C_p , its closeness centrality is not defined, because it has out-degree 0.

A vertex $y \in Y$ reaches $|C|$ vertices in 1 step, and hence, its closeness centrality is $\frac{|C|^2}{|C|(n-1)} = \frac{|C|}{n-1}$.

A vertex in C_i reaches $p - i$ other vertices, and their distance is $1, \dots, p - i$: consequently, its closeness centrality is $\frac{(p-i)^2}{\frac{(p-i)(p-i+1)}{2}(n-1)} = \frac{2(p-i)}{(n-1)(p-i+1)} \leq \frac{2}{n-1}$.

Finally, for a vertex $x \in X_1$ contained in N_x sets, for each $1 \leq i \leq p$, x reaches N_x vertices in C_i , and these vertices are at distance i . Hence, the closeness of x is $\frac{(pN_x)^2}{\frac{p(p+1)}{2}N_x(n-1)} = \frac{2pN_x}{(n-1)(p+1)} \leq \frac{2N_x}{n-1} \leq \frac{2|C|}{n-1}$. This concludes the proof. \square

Let us now compute the closeness centrality of a vertex $C \in C_0$. The reachable vertices are:

- all q vertices in Y , at distance 1;
- all $|C|q$ vertices in Z , at distance 2;
- $|X|$ vertices in X_1 or X_2 , at distance 1;
- R_C vertices in C_i for each i , at distance $i + 1$ (the sum of the distances of these vertices is $\sum_{i=1}^p i + 1 = -1 + \sum_{i=1}^{p+1} i = \frac{(p+2)(p+1)}{2} - 1$).

Hence, the closeness centrality of C is:

$$\begin{aligned} c(R_C) &= \frac{(q(1 + |C|) + |X| + pR_C)^2}{\left(q(1 + 2|C|) + |X| + \left(\frac{(p+1)(p+2)}{2} - 1\right) R_C\right) (n - 1)} \\ &= \frac{(q(1 + |C|) + |X| + pR_C)^2}{(q(1 + 2|C|) + |X| + g(p)R_C) (n - 1)}, \end{aligned}$$

where $g(p) = \frac{(p+1)(p+2)}{2} - 1$. We want to choose p and q verifying:

- (1) the closeness of vertices in C_0 is bigger than $\frac{2|C|}{n-1}$ (and hence bigger than the closeness of all other vertices);
- (2) $c(R_C)$ is a decreasing function of R_C for $0 \leq R_C \leq |C|$.

In order to satisfy Condition 2, the derivative $c'(R_C)$ of c is $(q(1 + |C|) + |X| + pR_C) \frac{[pg(p)R_C + 2p(q(1 + 2|C|) + |X|) - g(p)(q(1 + |C|) + |X|)]}{(q(1 + 2|C|) + |X| + g(p)R_C)^2(n - 1)}$.

This latter value is negative if and only if $pg(p)R_C + 2p(q(1 + 2|C|) + |X|) - g(p)(q(1 + |C|) + |X|) < 0$. Assuming $g(p) \geq 5p$ and $R_C < |C|$, this value is as follows:

$$\begin{aligned} &pg(p)R_C + 2p(q(1 + 2|C|) + |X|) - g(p)(q(1 + |C|) + |X|) \\ &\leq pg(p)|C| + 2pq + 4pq|C| + 2p|X| - g(p)(q - |C| - |X|) \\ &\leq pg(p)|C| + 4pq|C| - g(p)q|C| \\ &\leq pg(p)|C| - pq|C|. \end{aligned}$$

Assuming $q > g(p)$, we conclude that $c'(R_C) < 0$ for $0 \leq R_C \leq |C|$, and we verify Condition 2. In order to verify Condition 1, we want $c(R_C) \geq \frac{2|C|}{n+1}$ (since $c(R_C)$ is decreasing, it is enough $c(|C|) \geq \frac{2|C|}{n+1}$). Under the assumptions $q > g(p)$, $0 < |X| \leq |C|$ (which trivially holds for $|C|$ big enough, because $|X| \leq \log^p |C|$),

$$\begin{aligned} c(|C|) &= \frac{(q(1 + |C|) + |X| + pR_C)^2}{(q(1 + 2|C|) + |X| + g(p)R_C)(n - 1)} \\ &\geq \frac{q^2|C|^2}{(q(3|C|) + |C| + |C|)(n - 1)} \\ &\geq \frac{q|C|}{5(n - 1)} > \frac{2|C|}{n - 1}, \end{aligned}$$

if $q > 10$. To fulfill all required conditions, it is enough to choose $p = 7$, $g(p) = 35$, and $q = 36$.

B COMPARISON WITH THE STATE OF THE ART: DETAILED RESULTS

Table 6. Detailed Comparison of Running Times in Seconds, with $k = 1$

Directed Street								
Network	Nodes	Edges	OLH	OCL	DEGCut	DEGBOUND	NBCut	NBBOUND
faroe-islands	31,097	63,360	86.37	77.98	10.43	0.14	10.60	0.15
liechtenstein	54,972	112,773	738.75	408.83	71.71	1.40	72.00	1.44
isle-of-man	61,082	125,754	901.63	497.08	70.63	0.76	65.28	0.78
malta	91,188	190,768	1,362.68	814.23	110.28	3.97	122.80	3.91
belize	96,977	204,032	2,501.23	1,240.23	173.75	1.31	167.48	1.28
azores	237,174	474,050	2,949.90	692.69	150.96	6.76	154.73	6.75
Undirected Street								
Network	Nodes	Edges	OLH	OCL	DEGCut	DEGBOUND	NBCut	NBBOUND
faroe-islands	31,097	31,974	50.50	58.05	11.66	0.16	12.31	0.23
liechtenstein	54,972	56,616	282.06	313.99	84.21	1.19	82.48	1.46
isle-of-man	61,082	63,793	340.58	328.74	79.62	0.75	84.97	0.92
malta	91,188	101,437	504.75	519.12	139.15	2.60	140.52	2.60
belize	96,977	103,193	1,207.32	758.24	208.73	1.20	200.75	1.48
azores	237,174	243,185	1,344.92	484.92	192.05	3.89	177.38	5.08
Directed Complex								
Network	Nodes	Edges	OLH	OCL	DEGCut	DEGBOUND	NBCut	NBBOUND
polblogs	1,224	19,022	0.37	0.25	0.02	0.06	0.02	0.02
out.opsahl-openflights	2,939	30,501	1.15	2.43	0.03	0.21	0.03	0.03
ca-GrQc	5,241	28,968	1.85	3.32	0.09	0.28	0.09	0.20
out.subelj_jung-j_jung-j	6,120	50,535	1.03	<0.01	0.09	0.19	0.09	0.11
p2p-Gnutella08	6,301	20,777	1.95	0.73	0.04	0.34	0.03	0.28
out.subelj_jdk_jdk	6,434	53,892	1.21	<0.01	0.10	0.20	0.10	0.13
wiki-Vote	7,115	103,689	2.43	1.01	0.07	0.97	0.03	0.01
p2p-Gnutella09	8,114	26,013	3.00	1.15	0.11	0.55	0.10	0.50
ca-HepTh	9,877	51,946	7.67	15.75	0.27	1.29	0.26	0.95
freeassoc	10,617	72,172	3.17	7.54	1.30	3.09	1.29	3.15
ca-HepPh	12,006	236,978	20.49	47.02	6.72	3.20	6.68	2.62
out.lasagne-spanishbook	12,643	57,451	14.08	25.02	0.02	0.46	0.01	0.01
out.cfindex-google	15,763	170,335	19.09	29.60	1.53	5.03	2.33	4.50
ca-CondMat	23,133	186,878	54.34	118.92	0.24	5.03	0.19	1.08
out.subelj_cor_a_cor_a	23,166	91,500	18.07	2.79	3.43	6.50	3.39	4.90
out.ego-twitter	23,370	33,101	2.59	<0.01	0.02	0.08	0.02	0.02
out.ego_gplus	23,628	39,242	2.68	<0.01	0.02	0.04	0.02	0.02
as-caida20071105	26,475	106,762	45.78	116.88	0.05	16.84	0.04	0.05
cit-HepTh	27,769	352,768	69.01	16.18	14.60	10.21	14.41	8.42
Undirected Complex								
Network	Nodes	Edges	OLH	OCL	DEGCut	DEGBOUND	NBCut	NBBOUND
HC-BIOGRID	4,039	10,321	0.94	5.00	0.13	0.32	0.14	0.30
facebook_combined	4,039	88,234	2.08	7.80	0.35	0.12	0.34	0.02
Mus_musculus	4,610	5,747	0.79	2.30	0.02	0.27	0.01	0.13
Caenorhabditis_elegans	4,723	9,842	1.09	4.44	0.03	0.63	0.02	0.15
ca-GrQc	5,241	14,484	1.12	3.69	0.10	0.33	0.09	0.22
as20000102	6,474	12,572	0.77	5.97	0.01	0.64	<0.01	0.01
advogato	7,418	42,892	1.76	12.68	0.06	0.19	0.05	0.03
p2p-Gnutella09	8,114	26,013	3.68	18.02	0.13	2.79	0.11	0.35
hprd_pp	9,465	37,039	7.67	27.02	0.05	3.26	0.05	0.11
ca-HepTh	9,877	25,973	2.71	17.62	0.32	1.76	0.29	1.19
Drosophila_melanogaster	10,625	40,781	5.19	35.32	0.32	4.73	0.29	1.26
oregon1_010526	11,174	23,409	1.94	17.49	0.02	2.28	0.01	0.02
oregon2_010526	11,461	32,730	4.83	19.06	0.03	2.60	0.02	0.02
Homo_sapiens	13,690	61,130	7.12	60.31	0.04	7.87	0.03	0.12
GoogleNw	15,763	148,585	7.99	58.76	1.06	0.01	0.01	0.01
dip20090126_MAX	19,928	41,202	24.17	102.73	7.17	4.05	7.32	3.59
com-amazon.all.cmt	134,386	99,433	42.84	6.49	0.30	20.60	0.19	0.44

The running times refer to a sequential implementation of the algorithms.
Bold values represent better lower values.

Table 7. Detailed Comparison of Running Times in Seconds, with $k = 10$

Directed Street								
Network	Nodes	Edges	OLH	OCL	DEGCut	DEGBOUND	NBCut	NBBOUND
faroe-islands	31,097	63,360	93.32	77.33	10.68	0.17	10.75	0.17
liechtenstein	54,972	112,773	779.63	392.79	71.91	1.61	72.13	1.65
isle-of-man	61,082	125,754	904.73	473.63	66.54	0.82	66.47	0.84
malta	91,188	190,768	1,223.26	770.87	108.06	4.01	105.35	3.98
belize	96,977	204,032	2,457.49	1,083.29	167.77	1.35	165.42	1.33
azores	237,174	474,050	2,956.34	636.33	147.35	6.76	139.39	6.84
Undirected Street								
Network	Nodes	Edges	OLH	OCL	DEGCut	DEGBOUND	NBCut	NBBOUND
faroe-islands	31,097	31,974	52.02	57.75	12.31	0.23	11.85	0.18
liechtenstein	54,972	56,616	292.76	300.10	82.48	1.46	75.61	1.39
isle-of-man	61,082	63,793	337.80	330.02	84.97	0.92	83.93	0.84
malta	91,188	101,437	523.27	518.31	140.52	2.60	141.74	2.84
belize	96,977	103,193	1,083.27	759.01	200.75	1.48	200.55	1.36
azores	237,174	243,185	1,430.12	498.79	177.38	5.08	164.87	4.24
Directed Complex								
Network	Nodes	Edges	OLH	OCL	DEGCut	DEGBOUND	NBCut	NBBOUND
polblogs	1,224	19,022	0.36	0.25	0.04	0.07	0.04	0.04
out.opsahl-openflights	2,939	30,501	1.15	2.37	0.07	0.28	0.07	0.08
ca-GrQc	5,241	28,968	1.92	3.28	0.14	0.34	0.14	0.28
out.subelj_jung-j_jung-j	6,120	50,535	1.02	<0.01	0.09	0.18	0.09	0.11
p2p-Gnutella08	6,301	20,777	2.00	0.70	0.22	0.49	0.22	0.50
out.subelj_jdk_jdk	6,434	53,892	1.14	<0.01	0.10	0.20	0.10	0.13
wiki-Vote	7,115	103,689	2.31	1.00	0.11	0.98	0.11	0.08
p2p-Gnutella09	8,114	26,013	2.97	1.12	0.28	0.72	0.28	0.71
ca-HepTh	9,877	51,946	7.80	15.64	0.35	1.45	0.34	1.14
freeasoc	10,617	72,172	3.06	7.57	1.56	3.13	1.55	3.17
ca-HepPh	12,006	236,978	20.05	46.73	6.81	3.52	6.79	2.86
out.lasagne-spanishbook	12,643	57,451	13.86	24.59	0.06	6.45	0.05	0.09
out.cfindex-google	15,763	170,335	19.84	29.38	1.60	5.01	2.43	4.51
ca-CondMat	23,133	186,878	53.95	118.80	1.07	8.07	1.05	4.43
out.subelj_cor_a_cor_a	23,166	91,500	18.53	2.73	3.82	6.55	3.75	5.46
out.ego-twitter	23,370	33,101	2.64	<0.01	0.02	0.09	0.02	0.02
out.ego-gplus	23,628	39,242	2.75	<0.01	0.03	0.04	0.03	0.03
as-caida20071105	26,475	106,762	44.92	115.79	0.08	20.92	0.06	0.08
cit-HepTh	27,769	352,768	69.09	16.09	20.79	10.69	20.78	9.56
Undirected Complex								
Network	Nodes	Edges	OLH	OCL	DEGCut	DEGBOUND	NBCut	NBBOUND
HC-BIOGRID	4,039	10,321	0.93	5.16	0.14	0.30	0.20	0.39
facebook_combined	4,039	88,234	2.12	7.99	0.34	0.02	2.13	0.93
Mus_musculus	4,610	5,747	0.76	2.38	0.01	0.13	0.07	0.40
Caenorhabditis_elegans	4,723	9,842	1.11	4.34	0.02	0.15	0.04	0.71
ca-GrQc	5,241	14,484	1.12	3.61	0.09	0.22	0.15	0.39
as20000102	6,474	12,572	0.79	6.53	<0.01	0.01	0.02	1.57
advogato	7,418	42,892	2.91	12.84	0.05	0.03	0.13	1.81
p2p-Gnutella09	8,114	26,013	5.09	18.33	0.11	0.35	0.14	2.93
hprd_pp	9,465	37,039	6.86	27.53	0.05	0.11	0.38	3.70
ca-HepTh	9,877	25,973	2.81	17.62	0.29	1.19	0.41	2.00
Drosophila_melanogaster	10,625	40,781	4.95	35.50	0.29	1.26	0.44	5.27
oregon1_010526	11,174	23,409	4.47	20.29	0.01	0.02	0.04	4.83
oregon2_010526	11,461	32,730	4.79	22.56	0.02	0.02	0.06	5.31
Homo_sapiens	13,690	61,130	7.18	60.83	0.03	0.12	0.54	9.27
GoogleNw	15,763	148,585	7.92	79.34	0.01	0.01	4.82	15.20
dip20090126_MAX	19,928	41,202	22.33	103.93	7.32	3.59	10.33	5.79
com-amazon.all.cmt	134,386	99,433	42.38	6.58	0.19	0.44	0.37	20.98

The running times refer to a sequential implementation of the algorithms.
 Bold values represent better lower values.

Table 8. Detailed Comparison of Running Times in Seconds, with $k = 100$

Directed Street								
Network	Nodes	Edges	OLH	OCL	DEGCut	DEGBOUND	NBCut	NBBOUND
faroe-islands	31,097	63,360	91.17	77.98	11.59	0.31	11.69	0.32
liechtenstein	54,972	112,773	756.41	414.21	71.20	2.00	65.09	2.04
isle-of-man	61,082	125,754	914.43	473.11	70.98	1.32	70.27	1.34
malta	91,188	190,768	1,225.14	768.10	109.46	5.27	109.14	5.22
belize	96,977	204,032	2,441.20	1,179.81	171.79	2.00	170.13	1.98
azores	237,174	474,050	2,991.40	692.23	143.98	7.75	141.77	7.79
Undirected Street								
Network	Nodes	Edges	OLH	OCL	DEGCut	DEGBOUND	NBCut	NBBOUND
faroe-islands	31,097	31,974	56.98	57.84	11.85	0.18	12.53	0.24
liechtenstein	54,972	56,616	274.56	298.03	75.61	1.39	75.31	1.70
isle-of-man	61,082	63,793	327.95	335.44	83.93	0.84	86.36	1.01
malta	91,188	101,437	508.66	522.51	141.74	2.84	142.45	2.85
belize	96,977	103,193	1,100.30	768.48	200.55	1.36	203.28	1.52
azores	237,174	243,185	1,358.70	480.59	164.87	4.24	168.33	5.20
Directed Complex								
Network	Nodes	Edges	OLH	OCL	DEGCut	DEGBOUND	NBCut	NBBOUND
polblogs	1,224	19,022	0.38	0.25	0.14	0.08	0.14	0.06
out.opsahl-openflights	2,939	30,501	1.16	2.41	0.38	0.38	0.38	0.32
ca-GrQc	5,241	28,968	1.85	3.33	0.31	0.46	0.33	0.43
out.subelj_jung-j_jung-j	6,120	50,535	1.04	<0.01	0.09	0.19	0.09	0.11
p2p-Gnutella08	6,301	20,777	1.92	0.71	0.44	0.55	0.38	0.55
out.subelj_jdk_jdk	6,434	53,892	1.12	<0.01	0.11	0.21	0.11	0.13
wiki-Vote	7,115	103,689	2.50	1.01	0.30	1.20	0.31	0.19
p2p-Gnutella09	8,114	26,013	2.98	1.15	0.76	0.85	0.62	0.85
ca-HepTh	9,877	51,946	7.88	15.79	0.70	1.79	0.68	1.57
freeassoc	10,617	72,172	2.96	7.52	2.70	3.20	2.68	3.25
ca-HepPh	12,006	236,978	21.09	47.79	7.31	4.23	7.31	3.43
out.lasagne-spanishbook	12,643	57,451	14.43	25.02	0.29	6.57	0.20	0.25
out.cfinder-google	15,763	170,335	19.72	29.65	4.55	5.03	5.02	4.97
ca-CondMat	23,133	186,878	56.87	118.98	5.79	11.00	5.77	10.79
out.subelj_cora_cora	23,166	91,500	19.22	2.79	4.39	6.67	4.30	6.04
out.ego-twitter	23,370	33,101	2.77	<0.01	0.02	0.09	0.02	0.03
out.ego-gplus	23,628	39,242	2.83	<0.01	0.04	0.05	0.04	0.05
as-caida20071105	26,475	106,762	45.83	114.57	3.03	23.32	2.88	5.60
cit-HepTh	27,769	352,768	69.84	16.09	26.76	10.98	26.10	10.15
Undirected Complex								
Network	Nodes	Edges	OLH	OCL	DEGCut	DEGBOUND	NBCut	NBBOUND
HC-BIOGRID	4,039	10,321	0.94	5.05	0.20	0.39	0.21	0.39
facebook_combined	4,039	88,234	2.10	7.92	2.13	0.93	2.07	0.77
Mus_musculus	4,610	5,747	0.78	2.31	0.07	0.40	0.06	0.32
Caenorhabditis_elegans	4,723	9,842	1.08	4.50	0.04	0.71	0.04	0.19
ca-GrQc	5,241	14,484	1.11	3.71	0.15	0.39	0.14	0.29
as20000102	6,474	12,572	0.79	6.97	0.02	1.57	0.02	0.04
advogato	7,418	42,892	1.80	12.98	0.13	1.81	0.13	0.34
p2p-Gnutella09	8,114	26,013	3.75	18.33	0.14	2.93	0.13	0.37
hprd_pp	9,465	37,039	7.98	27.35	0.38	3.70	0.38	0.81
ca-HepTh	9,877	25,973	2.74	17.76	0.41	2.00	0.39	1.42
Drosophila_melanogaster	10,625	40,781	4.99	35.97	0.44	5.27	0.41	1.52
oregon1_010526	11,174	23,409	4.68	21.04	0.04	4.83	0.03	0.06
oregon2_010526	11,461	32,730	4.72	24.98	0.06	5.31	0.05	0.10
Homo_sapiens	13,690	61,130	7.25	60.75	0.54	9.27	0.54	1.00
GoogleNw	15,763	148,585	7.95	78.00	4.82	15.20	4.87	8.06
dip20090126_MAX	19,928	41,202	22.63	101.54	10.33	5.79	10.30	5.86
com-amazon.all.cmt	134,386	99,433	42.36	6.85	0.37	20.98	0.27	0.72

The running times refer to a sequential implementation of the algorithms.
Bold values represent better lower values.

Table 9. Detailed Comparison of the Edge Traversal Ratios, with $k = 1$

Directed Street						
Network	OLH	OCL	DEGCut	DEGBOUND	NBCut	NBBOUND
faroe-islands	4.080	3.742	4.125	338.011	4.086	437.986
liechtenstein	2.318	2.075	2.114	130.575	2.115	137.087
isle-of-man	2.623	3.740	2.781	224.566	2.769	314.856
malta	5.332	4.351	4.147	73.836	4.141	110.665
belize	2.691	3.969	2.606	253.866	2.595	444.849
azores	13.559	3.038	19.183	230.939	19.164	266.488
Undirected Street						
Network	OLH	OCL	DEGCut	DEGBOUND	NBCut	NBBOUND
faroe-islands	4.126	3.276	4.118	361.593	3.918	444.243
liechtenstein	2.318	2.027	2.107	171.252	2.122	183.240
isle-of-man	2.613	3.661	2.767	266.734	2.676	370.194
malta	4.770	4.164	3.977	122.729	3.958	232.622
belize	2.565	3.945	2.510	340.270	2.481	613.778
azores	22.406	2.824	18.654	589.985	18.810	727.528
Directed Complex						
Network	OLH	OCL	DEGCut	DEGBOUND	NBCut	NBBOUND
polblogs	3.201	1.131	31.776	1.852	31.974	5.165
out.opsahl-openflights	13.739	1.431	73.190	2.660	73.888	18.255
ca-GrQc	9.863	1.792	36.673	3.630	38.544	6.307
out.subelj_jung-j_jung-j	125.219	1.203	79.559	1.024	79.882	1.897
p2p-Gnutella08	5.696	1.121	66.011	4.583	81.731	6.849
out.subelj_jdk_jdk	116.601	1.167	74.300	1.023	74.527	1.740
wiki-Vote	9.817	2.760	261.242	1.479	749.428	395.278
p2p-Gnutella09	5.534	1.135	41.214	4.650	43.236	6.101
ca-HepTh	7.772	2.121	40.068	3.349	42.988	5.217
freeassoc	33.616	1.099	12.638	2.237	12.700	2.199
ca-HepPh	7.682	2.836	10.497	3.331	10.516	4.387
out.lasagne-spanishbook	13.065	2.553	1,871.296	7.598	6,786.506	3,160.750
out.cfindex-google	16.725	1.782	38.321	2.665	25.856	3.020
ca-CondMat	7.382	3.526	409.772	5.448	517.836	29.282
out.subelj_cora_cora	14.118	1.700	14.098	1.345	14.226	2.299
out.ego-twitter	2,824.713	1.000	1,870.601	28.995	3,269.183	278.214
out.ego-gplus	722.024	1.020	3,481.943	236.280	3,381.029	875.111
as-caida20071105	20.974	3.211	2,615.115	1.737	2,837.853	802.273
cit-HepTh	4.294	3.045	16.259	1.514	16.398	3.290
Undirected Complex						
Network	OLH	OCL	DEGCut	DEGBOUND	NBCut	NBBOUND
HC-BIOGRID	5.528	1.581	15.954	3.821	14.908	3.925
facebook_combined	10.456	3.726	56.284	18.786	56.517	98.512
Mus_musculus	18.246	1.743	70.301	3.253	104.008	7.935
Caenorhabditis_elegans	11.446	2.258	86.577	2.140	110.677	9.171
ca-GrQc	6.567	1.904	38.279	3.551	41.046	6.824
as20000102	19.185	2.402	1,550.351	3.213	1,925.916	498.000
advogato	8.520	2.018	315.024	18.181	323.163	142.654
p2p-Gnutella09	3.744	2.336	90.252	1.708	100.427	13.846
hprd_pp	6.543	2.397	392.853	2.091	407.261	63.953
ca-HepTh	7.655	2.075	42.267	3.308	46.326	5.593
Drosophila_melanogaster	5.573	2.346	69.457	1.822	75.456	6.904
oregon1_010526	20.474	3.723	1,603.739	2.703	1,798.822	399.071
oregon2_010526	17.330	4.748	1,138.475	2.646	1,227.105	520.955
Homo_sapiens	6.689	2.700	1,475.113	1.898	1,696.909	130.381
GoogleNw	15.591	8.389	107.902	15,763.000	15,763.000	15,763.000
dip20090126_MAX	2.883	3.826	5.833	6.590	5.708	7.392
com-amazon.all.cmt	415.286	2.499	5,471.982	3.297	8,224.693	373.294

Table 10. Detailed Comparison of the Edge Traversal Ratios, with $k = 10$

Directed Street						
Network	OLH	OCL	DEGCut	DEGBound	NBCut	NBBound
faroe-islands	3.713	2.884	4.037	290.626	4.025	361.593
liechtenstein	2.318	2.002	2.104	111.959	2.106	116.713
isle-of-man	2.623	2.933	2.711	209.904	2.720	288.123
malta	5.325	3.861	4.094	70.037	4.086	101.546
belize	2.690	3.638	2.592	244.275	2.580	416.210
azores	13.436	2.644	19.043	222.073	19.045	254.206
Undirected Street						
Network	OLH	OCL	DEGCut	DEGBound	NBCut	NBBound
faroe-islands	3.702	2.594	4.046	320.588	3.848	388.713
liechtenstein	2.316	1.965	2.097	142.047	2.114	150.608
isle-of-man	2.612	2.889	2.695	241.431	2.636	323.185
malta	4.768	3.615	3.920	115.574	3.910	208.192
belize	2.564	3.634	2.496	323.257	2.469	563.820
azores	22.392	2.559	18.541	539.032	18.712	653.372
Directed Complex						
Network	OLH	OCL	DEGCut	DEGBound	NBCut	NBBound
polblogs	3.199	1.039	13.518	1.496	13.544	2.928
out.opsahl-openflights	13.739	1.130	32.297	1.984	32.405	6.867
ca-GrQc	9.863	1.356	25.238	3.096	25.786	4.565
out.subelj_jung-j_jung-j	124.575	1.000	79.284	1.024	79.657	1.884
p2p-Gnutella08	5.684	1.064	12.670	3.241	12.763	3.599
out.subelj_jdk_jdk	116.228	1.000	74.106	1.023	74.363	1.730
wiki-Vote	9.812	1.205	166.941	1.453	174.775	25.411
p2p-Gnutella09	5.532	1.084	16.293	3.624	16.265	4.213
ca-HepTh	7.772	1.586	31.314	3.013	32.604	4.356
freeassoc	33.414	1.034	10.612	2.210	10.704	2.178
ca-HepPh	7.682	2.077	10.322	3.042	10.340	4.010
out.lasagne-spanishbook	13.063	1.483	303.044	1.067	351.262	94.351
out.cfinder-google	16.725	1.413	36.364	2.665	24.765	3.017
ca-CondMat	7.382	2.318	91.209	3.507	93.548	7.027
out.subelj_cora_cora	13.699	1.287	12.763	1.334	12.909	2.072
out.ego-twitter	2,689.884	1.000	1,817.032	28.157	2,872.213	218.411
out.ego-gplus	722.024	1.000	951.983	201.949	1,085.361	482.204
as-caida20071105	20.974	1.615	997.996	1.371	1,266.443	448.729
cit-HepTh	4.030	2.179	11.361	1.486	11.423	2.832
Undirected Complex						
Network	OLH	OCL	DEGCut	DEGBound	NBCut	NBBound
HC-BIOGRID	5.528	1.240	10.714	3.102	10.036	3.058
facebook_combined	10.456	1.292	9.103	2.236	9.371	2.694
Mus_musculus	18.246	1.316	18.630	2.279	20.720	3.288
Caenorhabditis_elegans	11.445	1.405	58.729	1.904	68.905	7.605
ca-GrQc	6.567	1.340	26.050	3.052	26.769	5.011
as20000102	19.185	1.529	196.538	1.314	209.674	52.210
advogato	8.520	1.405	131.173	2.043	132.207	11.155
p2p-Gnutella09	3.744	1.632	79.093	1.623	87.357	12.941
hprd_pp	6.543	1.436	47.945	1.837	47.866	8.620
ca-HepTh	7.655	1.546	32.612	2.961	34.407	4.677
Drosophila_melanogaster	5.573	1.672	50.840	1.646	54.637	5.743
oregon1_010526	20.474	1.451	418.099	1.282	429.161	109.549
oregon2_010526	17.330	1.560	364.277	1.302	371.929	71.186
Homo_sapiens	6.689	1.599	81.496	1.620	82.250	15.228
GoogleNw	15.591	1.320	23.486	1.252	23.053	2.420
dip20090126_MAX	2.881	1.836	4.055	4.556	4.065	4.498
com-amazon.all.cmt	414.765	1.618	3,407.016	3.279	3,952.370	199.386

Table 11. Detailed Comparison of the Edge Traversal Ratios, with $k = 100$

Directed Street						
Network	OLH	OCL	DEGCut	DEGBound	NBCut	NBBound
faroe-islands	3.713	2.823	3.694	150.956	3.691	168.092
liechtenstein	2.318	1.998	2.078	84.184	2.086	86.028
isle-of-man	2.620	2.902	2.551	139.139	2.567	167.808
malta	5.282	3.850	3.933	56.921	3.942	76.372
belize	2.688	3.617	2.526	184.718	2.516	268.634
azores	13.334	2.628	18.380	194.724	18.605	220.013
Undirected Street						
Network	OLH	OCL	DEGCut	DEGBound	NBCut	NBBound
faroe-islands	3.702	2.548	3.693	159.472	3.523	171.807
liechtenstein	2.311	1.959	2.072	96.782	2.095	99.768
isle-of-man	2.607	2.847	2.533	153.859	2.468	183.982
malta	4.758	3.605	3.745	89.929	3.730	137.538
belize	2.562	3.629	2.428	226.582	2.406	323.257
azores	22.345	2.548	18.092	411.760	18.384	476.253
Directed Complex						
Network	OLH	OCL	DEGCut	DEGBound	NBCut	NBBound
polblogs	3.198	1.037	3.951	1.245	3.961	1.731
out.opsahl-openflights	13.739	1.124	5.524	1.456	5.553	1.740
ca-GrQc	9.863	1.339	11.147	2.353	10.407	2.926
out.subelj_jung-j_jung-j	123.393	1.000	78.473	1.021	78.798	1.787
p2p-Gnutella08	5.684	1.063	6.611	2.935	7.750	3.278
out.subelj_jdk_jdk	114.210	1.000	73.522	1.020	73.755	1.669
wiki-Vote	9.812	1.186	61.375	1.236	60.475	9.436
p2p-Gnutella09	5.531	1.083	6.370	3.109	7.650	3.508
ca-HepTh	7.772	1.570	16.135	2.477	16.747	3.135
freeassoc	33.266	1.032	6.314	2.154	6.428	2.138
ca-HepPh	7.682	2.032	9.605	2.549	9.619	3.340
out.lasagne-spanishbook	13.063	1.467	56.689	1.043	80.069	33.271
out.cfinder-google	16.725	1.392	13.521	2.655	12.298	2.722
ca-CondMat	7.382	2.288	16.884	2.602	16.950	2.824
out.subelj_cora_cora	13.231	1.280	11.171	1.315	11.350	1.870
out.ego-twitter	2,621.659	1.000	1,574.836	26.893	1,908.731	110.236
out.ego-gplus	722.024	1.000	522.333	181.754	522.576	236.280
as-caida20071105	20.974	1.606	17.971	1.216	18.694	5.479
cit-HepTh	3.969	2.143	8.867	1.466	9.068	2.662
Undirected Complex						
Network	OLH	OCL	DEGCut	DEGBound	NBCut	NBBound
HC-BIOGRID	5.528	1.236	4.452	2.154	4.345	1.999
facebook_combined	10.456	1.292	3.083	1.470	3.074	1.472
Mus_musculus	18.245	1.305	7.940	1.944	9.518	2.631
Caenorhabditis_elegans	11.445	1.391	11.643	1.463	12.296	3.766
ca-GrQc	6.567	1.331	11.311	2.346	10.389	3.105
as20000102	19.185	1.512	7.318	1.174	7.956	3.593
advogato	8.520	1.398	32.629	1.706	33.166	7.784
p2p-Gnutella09	3.744	1.625	11.378	1.374	11.867	3.695
hprd_pp	6.543	1.422	21.053	1.547	22.191	3.468
ca-HepTh	7.655	1.539	16.406	2.454	17.030	3.301
Drosophila_melanogaster	5.573	1.655	29.115	1.487	30.979	4.614
oregon1_010526	20.474	1.443	13.300	1.163	14.611	6.569
oregon2_010526	17.330	1.530	18.203	1.173	21.758	7.258
Homo_sapiens	6.689	1.577	19.350	1.445	20.182	3.080
GoogleNw	15.591	1.320	16.224	1.172	16.506	2.010
dip20090126_MAX	2.880	1.815	2.789	2.602	2.784	2.546
com-amazon.all.cmt	414.765	1.605	1,368.675	3.236	1,654.150	97.735

C REAL-WORLD LARGE NETWORKS EXPERIMENTS: DETAILED RESULTS

Table 12. Results for Large Networks, Using DEGBOUND in Complex and NBBOUND in Street Networks

Directed Street								
Input	Nodes	Edges	$k = 1$		$k = 10$		$k = 100$	
			Impr.	Time	Impr.	Time	Impr.	Time
egypt	1,054,242	2,123,036	144.91	0:03:55	132.86	0:04:25	116.74	0:04:48
new_zealand	2,759,124	5,562,944	447.55	0:02:34	443.95	0:02:35	427.31	0:02:38
india	16,230,072	33,355,834	1,370.32	0:43:42	1,369.05	0:44:17	1,326.31	0:45:05
california	16,905,319	34,303,746	1,273.66	0:54:56	1,258.12	0:56:00	1,225.73	0:56:02
north_am	3,523,6615	70,979,433	1,992.68	2:25:58	1,967.87	2:29:25	1,877.78	2:37:14
Undirected Street								
Input	Nodes	Edges	$k = 1$		$k = 10$		$k = 100$	
			Impr.	Time	Impr.	Time	Impr.	Time
egypt	1,054,242	1,159,808	344.86	0:01:54	340.30	0:01:54	291.71	0:02:11
new_zealand	2,759,124	2,822,257	811.75	0:02:47	786.52	0:03:02	734.20	0:03:02
india	16,230,072	17,004,400	2,455.38	0:44:21	2,484.70	0:44:38	2,422.40	0:44:21
california	16,905,319	17,600,566	2,648.08	0:39:15	2,620.17	0:42:04	2,504.86	0:44:19
north_am	35,236,615	36,611,653	7,394.88	1:13:37	7,530.80	1:15:01	7,263.78	1:10:28
Directed Complex								
Input	Nodes	Edges	$k = 1$		$k = 10$		$k = 100$	
			Impr.	Time	Impr.	Time	Impr.	Time
cit-HepTh	27,769	352,768	16.34	0:00:01	11.41	0:00:01	9.06	0:00:02
cit-HepPh	34,546	421,534	23.68	0:00:01	19.88	0:00:01	14.41	0:00:02
p2p-Gnut31	62,586	147,892	194.19	0:00:01	44.24	0:00:01	19.34	0:00:04
soc-Eps1	75,879	508,837	243.14	0:00:01	43.75	0:00:01	33.60	0:00:05
soc-Slash0811	77,360	828,161	1,007.70	0:00:00	187.46	0:00:00	21.09	0:00:18
twitter_comb	81,306	2,684,592	1,024.32	0:00:01	692.96	0:00:01	145.68	0:00:05
Slash090221	82,140	549,202	177.82	0:00:02	162.30	0:00:02	108.53	0:00:03
gplus_comb	107,614	24,476,570	1,500.35	0:00:04	235.17	0:00:04	62.54	0:02:19
soc-sign-eps	131,828	840,799	225.91	0:00:03	161.58	0:00:03	39.26	0:00:16
email-EuAll	265,009	418,956	4,724.80	0:00:00	3,699.48	0:00:00	1,320.22	0:00:01
web-Stanford	281,903	2,312,497	13.59	0:04:00	8.70	0:04:00	7.47	0:07:15
web-NotreD	325,729	1,469,679	1,690.08	0:00:02	132.83	0:00:02	66.88	0:00:49
amazon0601	403,394	3,387,388	10.81	0:14:54	8.87	0:14:54	6.84	0:22:04
web-BerkStan	685,230	7,600,595	3.95	1:36:21	3.67	1:36:21	3.47	1:49:12
web-Google	875,713	5,105,039	228.61	0:01:51	96.63	0:01:51	38.69	0:10:29
youtube-links	1,138,494	4,942,297	662.78	0:01:33	200.68	0:01:33	125.72	0:07:02
in-2004	1,382,870	16,539,643	43.68	0:41:45	29.89	0:41:45	16.68	1:48:42
trec-wt10g	1,601,787	8,063,026	33.86	0:36:01	20.39	0:36:01	16.73	1:10:54
soc-pokec	1,632,803	22,301,964	21,956.64	0:00:17	2,580.43	0:06:14	1,106.90	0:12:35
zhishi-hudong	1,984,484	14,682,258	30.37	1:25:38	27.71	1:25:38	24.95	1:53:27
zhishi-baidu	2,141,300	17,632,190	44.05	1:17:52	38.61	1:17:52	23.17	3:08:05
wiki-Talk	2,394,385	5,021,410	34,863.42	0:00:08	28,905.76	0:00:08	9,887.18	0:00:18
cit-Patents	3,774,768	16,518,947	9,454.04	0:02:07	8,756.77	0:02:07	8,340.18	0:02:13

(Continued)

Table 12. Continued
Undirected Complex

Input	Nodes	Edges	$k = 1$		$k = 10$		$k = 100$	
			Impr.	Time	Impr.	Time	Impr.	Time
ca-HepPh	12,008	118,489	10.37	0:00:00	10.20	0:00:00	9.57	0:00:01
CA-AstroPh	18,772	198,050	62.47	0:00:00	28.87	0:00:01	14.54	0:00:01
CA-CondMat	23,133	93,439	247.35	0:00:00	84.48	0:00:00	17.06	0:00:01
email-Enron	36,692	183,831	365.92	0:00:00	269.80	0:00:00	41.95	0:00:01
loc-brightkite	58,228	214,078	308.03	0:00:00	93.85	0:00:01	53.49	0:00:02
flickrEdges	105,938	2,316,948	39.61	0:00:23	17.89	0:00:55	15.39	0:01:16
gowalla	196,591	950,327	2,412.26	0:00:01	33.40	0:01:18	28.13	0:01:33
com-dblp	317,080	1,049,866	500.83	0:00:10	300.61	0:00:17	99.64	0:00:52
com-amazon	334,863	925,872	37.76	0:02:21	31.33	0:02:43	18.68	0:04:34
com-lj.all	477,998	530,872	849.57	0:00:07	430.72	0:00:13	135.14	0:00:45
com-youtube	1,134,890	2,987,624	2,025.32	0:00:32	167.45	0:06:44	110.39	0:09:16
soc-pokec	1,632,803	30,622,564	46,725.71	0:00:18	8,664.33	0:02:16	581.52	0:18:12
as-skitter	1,696,415	11,095,298	185.91	0:19:06	164.24	0:21:53	132.38	0:27:06
com-orkut	3,072,441	117,185,083	23,736.30	0:02:32	255.17	2:54:58	69.23	15:02:06
youtube-u-g	3,223,585	9,375,374	11,473.14	0:01:07	91.17	2:07:23	66.23	2:54:12

D IMDB CASE STUDY: DETAILED RESULTS

Table 13. Detailed Ranking of the IMDB Actor Graph

	1940	1945	1950	1955
1	Semels, Harry (I)	Corrado, Gino	Flowers, Bess	Flowers, Bess
2	Corrado, Gino	Steers, Larry	Steers, Larry	Harris, Sam (II)
3	Steers, Larry	Flowers, Bess	Corrado, Gino	Steers, Larry
4	Bracey, Sidney	Semels, Harry (I)	Harris, Sam (II)	Corrado, Gino
5	Lucas, Wilfred	White, Leo (I)	Semels, Harry (I)	Miller, Harold (I)
6	White, Leo (I)	Mortimer, Edmund	Davis, George (I)	Farnum, Franklyn
7	Martell, Alphonse	Boteler, Wade	Magrill, George	Magrill, George
8	Conti, Albert (I)	Phelps, Lee (I)	Phelps, Lee (I)	Conaty, James
9	Flowers, Bess	Ring, Cyril	Ring, Cyril	Davis, George (I)
10	Sedan, Rolfe	Bracey, Sidney	Moorhouse, Bert	Cording, Harry
	1960	1965	1970	1975
1	Flowers, Bess	Flowers, Bess	Flowers, Bess	Flowers, Bess
2	Harris, Sam (II)	Harris, Sam (II)	Harris, Sam (II)	Harris, Sam (II)
3	Farnum, Franklyn	Farnum, Franklyn	Tamiroff, Akim	Tamiroff, Akim
4	Miller, Harold (I)	Miller, Harold (I)	Farnum, Franklyn	Welles, Orson
5	Chefe, Jack	Holmes, Stuart	Miller, Harold (I)	Sayre, Jeffrey
6	Holmes, Stuart	Sayre, Jeffrey	Sayre, Jeffrey	Miller, Harold (I)
7	Steers, Larry	Chefe, Jack	Quinn, Anthony (I)	Farnum, Franklyn
8	Paris, Manuel	Paris, Manuel	O'Brien, William H.	Kemp, Kenner G.
9	O'Brien, William H.	O'Brien, William H.	Holmes, Stuart	Quinn, Anthony (I)
10	Sayre, Jeffrey	Stevens, Bert (I)	Stevens, Bert (I)	O'Brien, William H.

(Continued)

Table 13. Continued

	1980	1985	1990	1995
1	Flowers, Bess	Welles, Orson	Welles, Orson	Lee, Christopher (I)
2	Harris, Sam (II)	Flowers, Bess	Carradine, John	Welles, Orson
3	Welles, Orson	Harris, Sam (II)	Flowers, Bess	Quinn, Anthony (I)
4	Sayre, Jeffrey	Quinn, Anthony (I)	Lee, Christopher (I)	Pleasence, Donald
5	Quinn, Anthony (I)	Sayre, Jeffrey	Harris, Sam (II)	Hitler, Adolf
6	Tamiroff, Akim	Carradine, John	Quinn, Anthony (I)	Carradine, John
7	Miller, Harold (I)	Kemp, Kenner G.	Pleasence, Donald	Flowers, Bess
8	Kemp, Kenner G.	Miller, Harold (I)	Sayre, Jeffrey	Mitchum, Robert
9	Farnum, Franklyn	Niven, David (I)	Tovey, Arthur	Harris, Sam (II)
10	Niven, David (I)	Tamiroff, Akim	Hitler, Adolf	Sayre, Jeffrey
	2000	2005	2010	2014
1	Lee, Christopher (I)	Hitler, Adolf	Hitler, Adolf	Madsen, Michael (I)
2	Hitler, Adolf	Lee, Christopher (I)	Lee, Christopher (I)	Trejo, Danny
3	Pleasence, Donald	Steiger, Rod	Hopper, Dennis	Hitler, Adolf
4	Welles, Orson	Sutherland, Donald (I)	Keitel, Harvey (I)	Roberts, Eric (I)
5	Quinn, Anthony (I)	Pleasence, Donald	Carradine, David	De Niro, Robert
6	Steiger, Rod	Hopper, Dennis	Sutherland, Donald (I)	Dafoe, Willem
7	Carradine, John	Keitel, Harvey (I)	Dafoe, Willem	Jackson, Samuel L.
8	Sutherland, Donald (I)	von Sydow, Max (I)	Caine, Michael (I)	Keitel, Harvey (I)
9	Mitchum, Robert	Caine, Michael (I)	Sheen, Martin	Carradine, David
10	Connery, Sean	Sheen, Martin	Kier, Udo	Lee, Christopher (I)

Table 14. Detailed Edge Traversal Ratios on the IMDB Actor Graph

YEAR	1940	1945	1950	1955
NODES	69,011	83,068	97,824	120,430
EDGES	3,417,144	5,160,584	6,793,184	8,674,159
IMPR ($k = 1$)	51.74	61.46	67.50	91.46
IMPR ($k = 10$)	32.95	40.73	44.72	61.52
YEAR	1960	1965	1970	1975
NODES	146,253	174,826	210,527	257,896
EDGES	11,197,509	12,649,114	14,209,908	16,080,065
IMPR ($k = 1$)	122.63	162.06	211.05	285.57
IMPR ($k = 10$)	80.50	111.51	159.32	221.07
YEAR	1980	1985	1990	1995
NODES	310,278	375,322	463,078	557,373
EDGES	18,252,462	20,970,510	24,573,288	28,542,684
IMPR ($k = 1$)	380.52	513.40	719.21	971.11
IMPR ($k = 10$)	296.24	416.27	546.77	694.72
YEAR	2000	2005	2010	2014
NODES	681,358	880,032	1,237,879	1,797,446
EDGES	33,564,142	41,079,259	53,625,608	72,880,156
IMPR ($k = 1$)	1,326.53	1,897.31	2,869.14	2,601.52
IMPR ($k = 10$)	838.53	991.89	976.63	1,390.32

REFERENCES

- Amir Abboud, Fabrizio Grandoni, and Virginia V. Williams. 2015. Subcubic equivalences between graph centrality problems, APSP and diameter. In *Proceedings of the 26th ACM/SIAM Symposium on Discrete Algorithms (SODA'15)*. 1681–1697. Retrieved from <http://people.idsia.ch/~grandoni/Pubblicazioni/AGV15soda.pdf>.
- Amir Abboud and Virginia V. Williams. 2014. Popular conjectures imply strong lower bounds for dynamic problems. In *Proceedings of the 55th Annual Symposium on Foundations of Computer Science (FOCS'14)*. 434–443. Retrieved from <http://arxiv.org/abs/1402.0054>.
- Amir Abboud, Virginia V. Williams, and Joshua Wang. 2016. Approximation and fixed parameter subquadratic algorithms for radius and diameter. In *Proceedings of the 27th ACM/SIAM Symposium on Discrete Algorithms (SODA'16)*. 377–391.
- Amir Abboud, Virginia V. Williams, and Oren Weimann. 2014. Consequences of faster alignment of sequences. In *Proceedings of the 41st International Colloquium on Automata, Languages and Programming (ICALP'14)*. 39–51. Retrieved from <http://www.cs.haifa.ac.il/~oren/Publications/hardstrings.pdf>.
- Alex Bavelas. 1950. Communication patterns in task-oriented groups. *Journal of the Acoustical Society of America* 22, 6 (1950), 725–730.
- Elisabetta Bergamini, Michele Borassi, Pierluigi Crescenzi, Andrea Marino, and Henning Meyerhenke. 2016. Computing top- k closeness centrality faster in unweighted graphs. In *Proceedings of the Meeting on Algorithm Engineering and Experiments (ALENEX'16)*.
- Paolo Boldi and Sebastiano Vigna. 2013. In-core computation of geometric centralities with hyperball: A hundred billion nodes and beyond. In *Proceedings of the 13th IEEE International Conference on Data Mining Workshops (ICDM'13)*. 621–628.
- Paolo Boldi and Sebastiano Vigna. 2014. Axioms for centrality. *Internet Mathematics* 10, 3–4 (2014), 222–262. Retrieved from <http://www.tandfonline.com/doi/abs/10.1080/15427951.2013.865686>.
- Michele Borassi. 2016. A note on the complexity of computing the number of reachable vertices in a digraph. *Information Processing Letters* 116, 10 (2016), 628–630.
- Michele Borassi, Pierluigi Crescenzi, and Michel Habib. 2015a. Into the square – On the complexity of some quadratic-time solvable problems. In *Proceedings of the 16th Italian Conference on Theoretical Computer Science (ICTCS'15)*. 1–17. Retrieved from <http://arxiv.org/abs/1407.4972>.
- Michele Borassi, Pierluigi Crescenzi, and Andrea Marino. 2015b. Fast and simple computation of top- k closeness centralities. arXiv:1507.01490.
- Ulrik Brandes and Daniel Fleischer. 2005. Centrality measures based on current flow. In *Proceedings of the 22nd Annual Symposium on Theoretical Aspects of Computer Science (STACS'05)*. Volker Diekert and Bruno Durand (Eds.), Lecture Notes in Computer Science, Vol. 3404, Springer, 533–544. DOI: https://doi.org/10.1007/978-3-540-31856-9_44
- Shiri Chechik, Edith Cohen, and Haim Kaplan. 2015. Average distance queries through weighted samples in graphs and metric spaces: High scalability with tight statistical guarantees. In *Proceedings of the Approximation, Randomization, and Combinatorial Optimization. Algorithms and Techniques (APPROX/RANDOM'15)*. LIPIcs, Vol. 40, Schloss Dagstuhl - Leibniz-Zentrum fuer Informatik, 659–679. DOI: <https://doi.org/10.4230/LIPIcs.APPROX-RANDOM.2015.659>
- Duanbing Chen, Linyuan Lu, Ming-Sheng Shang, Yi-Cheng Zhang, and Tao Zhou. 2012. Identifying influential nodes in complex networks. *Physica A: Statistical Mechanics and its Applications* 391, 4 (2012), 1777–1787. DOI: <https://doi.org/10.1016/j.physa.2011.09.017>
- Edith Cohen, Daniel Delling, Thomas Pajor, and Renato F. Werneck. 2014. Computing classic closeness centrality, at scale. In *Proceedings of the 2nd ACM Conference on Online Social Networks (COSN'14)*. 37–50.
- Thomas H. Cormen, Charles E. Leiserson, Ronald L. Rivest, and Clifford Stein. 2009. *Introduction to Algorithms* (3rd ed.). MIT Press.
- Gábor Csárdi and Tamás Nepusz. 2006. The igraph software package for complex network research. *InterJournal, Vol: Complex Systems* 1695, 5 (2006), 1–9. Retrieved from <http://www.necsi.edu/events/iccs6/papers/c1602a3c126ba822d0bc4293371c.pdf>.
- David Eppstein and Joseph Wang. 2004. Fast approximation of centrality. *Journal of Graph Algorithms and Applications* 8, 1 (2004), 39–45. Retrieved from <http://www.emis.ams.org/journals/JGAA/accepted/2004/EppsteinWang2004.8.1.pdf>.
- François Le Gall. 2012. Faster algorithms for rectangular matrix multiplication. In *Proceedings of the 53rd Annual IEEE Symposium on Foundations of Computer Science (FOCS'12)*. IEEE Computer Society, 514–523. DOI: <https://doi.org/10.1109/FOCS.2012.80>
- François Le Gall. 2014. Powers of tensors and fast matrix multiplication. In *Proceedings of the 2014 International Symposium on Symbolic and Algebraic Computation (ISSAC'14)*. 296–303. DOI: <https://doi.org/10.1145/2608628.2608664>
- Aric A. Hagberg, Daniel A. Schult, and Pieter J. Swart. 2008. Exploring network structure, dynamics, and function using NetworkX. In *Proceedings of the 7th Python in Science Conference (SCIPY'08)*. 11–15. Retrieved from https://conference.scipy.org/proceedings/scipy2008/paper_2/full_text.pdf.

- Russell Impagliazzo, Ramamohan Paturi, and Francis Zane. 2001. Which problems have strongly exponential complexity? *Journal of Computer and System Sciences* 63, 4 (Dec. 2001), 512–530. DOI : <https://doi.org/10.1006/jcss.2001.1774>
- U. Kang, Spiros Papadimitriou, Jimeng Sun, and Tong Hanghang. 2011. Centralities in large networks: Algorithms and observations. In *Proceedings of the SIAM International Conference on Data Mining (SDM'11)*. 119–130. Retrieved from <http://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.231.8735>.
- Erwan Le Merrer, Nicolas Le Scouarnec, and Gilles Trédan. 2014. Heuristical Top-k: Fast estimation of centralities in complex networks. *Information Processing Letters* 114, 8 (2014), 432–436.
- Yeon-Sup Lim, Daniel S. Menasché, Bruno Ribeiro, Don Towsley, and Prithwish Basu. 2011. Online estimating the k central nodes of a network. In *Proceedings of the 2011 IEEE Network Science Workshop (NS'11)*. 118–122.
- Nan Lin. 1976. *Foundations of Social Research*. McGraw-Hill. Retrieved from <http://books.google.it/books?id=DIowAAAAAAJ>.
- Massimo Marchiori and Vito Latora. 2000. Harmony in the small-world. *Physica A: Statistical Mechanics and Its Applications* 285, 3–4 (Oct. 2000), 539–546. Retrieved from <https://www.sciencedirect.com/science/article/pii/S0378437100003113>.
- Mark E. J. Newman. 2010. *Networks: An Introduction*. OUP Oxford. Retrieved from <http://books.google.it/books?id=q7HVtpYVfC0C>
- Kazuya Okamoto, Wei Chen, and X. Y. Li. 2008. Ranking of closeness centrality for large-scale social networks. *Frontiers in Algorithmics* 5059 (2008), 186–195. Retrieved from http://link.springer.com/chapter/10.1007/978-3-540-69311-6_21.
- Paul W. Olsen, Alan G. Labouseur, and Jeong-Hyon Hwang. 2014. Efficient top-k closeness centrality search. In *Proceedings of the 30th IEEE International Conference on Data Engineering (ICDE'14)*. 196–207. Retrieved from http://ieeexplore.ieee.org/xpls/abs_all.jsp?arnumber=6816651.
- Mihai Pătraşcu and Ryan Williams. 2010. On the possibility of faster SAT algorithms. In *Proceedings of the 21st ACM/SIAM Symposium on Discrete Algorithms (SODA'10)*. Retrieved from <http://epubs.siam.org/doi/abs/10.1137/1.9781611973075.86>.
- Liam Roditty and Virginia V. Williams. 2013. Fast approximation algorithms for the diameter and radius of sparse graphs. In *Proceedings of the 45th Annual ACM Symposium on Theory of Computing (STOC'13)*. 515–524. DOI : <https://doi.org/10.1145/2488608.2488673>
- Ahmet E. Sariyüce, Kamer Kaya, Erik Saule, and Umit V. Catalyurek. 2013. Incremental algorithms for closeness centrality. In *Proceedings of the IEEE International Conference on Big Data (ICBDA'13)*. 118–122.
- Jeremy G. Siek, Lie Quan Lee, and Andrew Lumsdaine. 2001. *The Boost Graph Library: User Guide and Reference Manual*. Pearson Education. Retrieved from <http://books.google.it/books?id=CPi7g1hjlYIC>.
- Christian L. Staudt, Aleksejs Sazonovs, and Henning Meyerhenke. 2016. NetworKit: A tool suite for large-scale complex network analysis. *Network Science* 4, 4 (2016), 508–530. DOI : <https://doi.org/10.1017/nws.2016.20>
- William Stein and David Joyner. 2005. Sage: System for algebra and geometry experimentation. *SIGSAM Bulletin* 39, 2 (2005), 61–64. Retrieved from <https://doi.org/10.1145/1101884.1101889>
- Wei Wang and Choon Yik Tang. 2014. Distributed computation of classic and exponential closeness on tree graphs. In *Proceedings of the American Control Conference (ACC'14)*. 2090–2095.
- Stanley Wasserman and Katherine Faust. 1994. *Social Network Analysis: Methods and Applications*. Cambridge University Press. Retrieved from <http://books.google.it/books?id=CAm2DplqRUIC>.
- Ryan Williams. 2005. A new algorithm for optimal 2-constraint satisfaction and its implications. *Theoretical Computer Science* 348, 2–3 (2005), 357–365. DOI : <https://doi.org/10.1016/j.tcs.2005.09.023>
- Virginia V. Williams. 2012. Multiplying matrices faster than coppersmith-winograd. In *Proceedings of the 44th Annual ACM Symposium on Theory of Computing (STOC'12)*. 887–898. DOI : <https://doi.org/10.1145/2213977.2214056>
- Virginia V. Williams and Ryan Williams. 2010. Subcubic equivalences between path, matrix and triangle problems. In *Proceedings of the 51st IEEE Annual Symposium on Foundations of Computer Science (FOCS'10)*. 645–654.

Received November 2016; revised November 2018; accepted June 2019